

ACRME — Technical Design Document (TDD)

Title	Azure Capacity Reservation Management Engine (ACRME) — Technical Design Document
Version	1.0 (net-new)
Date	2 September 2026
Status	Draft for review — supersedes the Production Readiness Review as the technical design of record
Baseline	Azure Capacity & Quota Management — Consolidated Requirements Baseline v2.4 (7 Sep 2026)
Owner	Vishnuvardhan Reddy · Principal Cloud Architect
Audience	Engineering, SRE/platform operations, security, POC leads
Companion	Functional Design Document (<code>acrme_functional_design_document.md</code>)
Decision records	ADR-001 (region selection), ADR-002 (quota & capacity, single-pool), ADR-003 (capacity during DR), ADR-004 (forecast & increase), ADR-005 (distributed DR reference model) — all reconciled to v2.4

Purpose. This TDD describes **how** ACRME is built — components, runtime, topology, data, state, algorithms, interfaces, security, observability, NFRs, integration, and POC gating — traceable to Requirements Baseline v2.4 and to the companion FDD (which owns the what). This document is **self-contained**: all normative algorithms, formulas, schemas, enumerations, and constants are inlined, not referenced externally.

Reconciliation note (v2.4 — reservation-model gaps). This revision folds the reviewed architecture-diagram gaps into the technical design. **Reservation eligibility** is refined: Availability-Set VMs are ineligible for zonal capacity reservations (**CAP-020**), and onboarding requires a **deallocate/redeploy-to-AZ precondition** for any VM not already zone-pinned (**CAP-021**). The managed estate is initialised as a **seed matrix of count-0 (seed) reservations** across every eligible SKU×region×AZ combination, governed by per-product-team budget (**CAP-022**, extending CAP-009), reconciled with **reactive SKU/AZ discovery** which auto-creates a reservation and simultaneously raises a scope-file governance item (**CAP-024**, reconciling CAP-019). Reservations are organised into an explicit **regional + per-AZ CRG structure per environment** (**CAP-023**, extending CAP-011 and refining the three-CRG model in Section 5.2). VM placement targets an **even ≈1/zone_count per-zone distribution with a rebalancing action** (**PLC-011**, Section 8.5). A shared **core subscription is classified entirely production** (**CAP-001a**). The **decommissioning-workflow boundary** is made explicit (**CAP-008/CAP-010**): automatic reconciliation only right-sizes a reservation down to `Allocated + Buffer` (never to zero or delete on its own), while reducing to zero, retirement and deletion are **gated** workflow actions. A deterministic **RG/CRG/subscription naming convention + counter** is adopted (**OPS-006**, config **C-12**); the zone-distribution tolerance is config **C-13**. A “**seed reservation**” (count-0 reservation) is distinct from the placement “**seed record**” (`CustomerSeedRecord` , PLC-003).

Reconciliation note (decisions carried into v2.4, first confirmed at v2.3). This TDD implements the following confirmed decisions, unchanged in v2.4: **single governed quota pool** as the primary technical model with logical earmarks for Prod/DR protection (QUA-004, ADR-002); **weighted cross-geo DR for the Middle East — [Amended v2.4]** Middle East Prod+CVAL co-locate in a weighted-selected Middle East Standard region and DR is placed cross-geo in a weighted-selected **Europe** Standard region (Switzerland North remains the default example, not a fixed region), both selections using the weighted capacity model (REG-002, DR-020, ADR-005); this supersedes the former `DR_NOT_OFFERED` Middle East position (DEC-001 **RESOLVED**); **max-not-sum** destination DR sizing with an authoritative `SourceDestinationDRIndex` (DR-016/017/018, ADR-005); **exact-production-region-first** validation with a governed `CustomerSeedRecord` (PLC-001..005, ADR-001); **five-state engine machine** and **standby activation waves** (DR-019, ADR-003).

Region model (v2.4, baseline Section 6). The catalogue now defines **five geographies** each carrying an explicit `distribution_model`: the **US is the only three-region geography** (West US 3, Central US, Canada Central; East US 2 Restricted); **EU** (Switzerland North + Sweden Central; North Europe & West Europe Restricted), **Australia** (Australia East + Australia Southeast), **Asia Pacific** (East Asia + Southeast Asia; Japan East **pending**) and the **Middle East** (UAE North + Saudi Arabia Central) are **two-region geographies**. In the EU, Australia and Asia Pacific two-region geographies, **CVAL and DR co-locate** in the non-production region (**PLC-010a**). **[Amended v2.4]** The **Middle East is now a cross-geo DR geography**: Prod and CVAL co-locate in a weighted-selected Middle East Standard region (separate CRGs), while **DR is placed cross-geo in a weighted-selected Europe Standard region** — both selections use the weighted capacity model (PLC-010b overrides PLC-010a for the Middle East; DR-020, REG-002/REG-003). The region catalogue and `distribution_model` are configuration-driven (REG-001); adding Japan East or changing the Europe DR destination set are config changes with no code change.

1. Introduction

1.1 Scope

This document specifies the technical design of ACRME: the control engine that guarantees every managed Azure deployment has **both** reserved physical capacity and deployable VM-family quota — in the correct region/zone/SKU — before it proceeds, while minimising idle cost through lean distributed DR and continuous reconciliation. It covers component architecture, runtime and deployment, data and state, the placement/scoring/DR algorithms, the API surface, security, observability, non-functional behaviour, integration, and the POC dependencies that gate production reliance.

1.2 Relationship to the FDD

The FDD (`acrme_functional_design_document.md`) defines **what** ACRME does — capabilities, flows, states, and rules — traceable to Baseline v2.4. This TDD defines **how** those capabilities are realised. Where the FDD names a behaviour (e.g. “single governed pool”, “max-not-sum DR”, “readiness states”), this TDD gives the component, schema, algorithm, and interface that implements it. Section references to the FDD use the form (FDD Section 4.x).

1.3 Decision records referenced

ADR	Title	Key technical mandate used here
ADR-001	Region selection & customer placement	Exact-region-first validation; seed record; staged placement pipeline; input modes
ADR-002	Quota & capacity management	Single governed pool (primary) ; logical earmarks; exception two-group topology
ADR-003	Capacity management during DR	Five-state engine machine; three-tier emergency transfer; quota-neutral Tier 2
ADR-004	Forecast & increase of capacity/quota	Reservation floor vs forecast growth; auto-increase triggers; 10-step lifecycle
ADR-005	Distributed DR reference model	<code>SourceDestinationDRIndex</code> ; max-not-sum; overcommit ratio/gap; Section 12A worked model

1.4 Evidence tags

Every non-trivial design assertion carries an evidence tag consistent with the requirements corpus:

[Documented] (Microsoft Learn / Azure platform), [Decided] (approved ADR decision), [Derived] (lo-

gical consequence; POC where noted), [Assumed] (tunable policy default, not yet empirically validated).

2. Architecture Overview

2.1 Design principles (baseline Section 20)

1. **Readiness is capacity $\wedge$ quota $\wedge$ freshness.** No deployment is “ready” on a reservation alone; deployable consumer quota and non-state state are equally required (RDY-003, NFR-002). [Decided]
2. **Continuous reconciliation of desired vs actual.** A reservation floor (`Allocated + Buffer` , CAP-003) is enforced every cycle; forecast growth is a separate proactive path (ADR-004). [Decided]
3. **Lean, distributed, reciprocal DR.** No dedicated DR regions; each region hosts Prod + CVAL + DR standby for other regions; standby is sized max-not-sum and shared across mutually exclusive failures (DR-016/017, ADR-005). [Decided]
4. **Single governed quota pool.** Collect all eligible per-region quota into one pool per region/quota family and allocate on demand; protect Prod/DR by logical earmark, minimising quota-increase requests to Microsoft (QUA-004, ADR-002). [Decided]
5. **Config-as-code, deterministic, replayable.** All policy (weights, buffers, thresholds, catalogue) is versioned; every decision records inputs, policy version, and snapshot ref for deterministic replay. [Decided]
6. **Fail-safe and idempotent.** Stale state blocks placement; all mutations are idempotent with operation handles and optimistic concurrency (NFR-002/007). [Decided]

2.2 Constraints (baseline Section 21)

- **[Amended v2.4] Middle East DR is offered cross-geo into Europe (DEC-001 RESOLVED).** The Middle East programme is legal-owned and serves government/medical customers under data-residency/sovereignty laws. Per the latest legal position, **DR is now provided** for the Middle East: Prod and CVAL are provisioned in a weighted-selected **Middle East Standard region** (separate CRGs per ENV-003), and **DR is placed cross-geo in a weighted-selected Europe Standard region** — Europe being an approved DR destination for the Middle East (assumption **A-ME1**; per-country carve-outs are handled via configuration and are out of scope of the engine). Both region selections use the **weighted capacity model** (DR-020, REG-002/REG-003, PLC-010b). The placement/DR engine emits a real `dr_region` in Europe for Middle East sources, creates a `SourceDestinationDRIndex` entry, reserves a DR earmark **at the Europe destination**, and performs cross-geo activation. Because CVAL now co-locates with **Prod locally** (not with DR), the CVAL-sacrifice DR bootstrap (DR-005/006) does **not** apply to the Middle East; Middle East DR uses **dedicated reserved capacity in Europe**. [Documented]
- **Azure Quota Group / `groupQuotas` maturity** gates single-pool enforcement (DEP-001); the two-group topology is the sanctioned fallback. [Documented]
- **No native intra-pool sub-reservation** in Azure quota groups — the DR earmark and Prod floor are engine-enforced arithmetic, not platform primitives (Scenario 9). [Documented]
- **ARM throttling budgets** bound reconciliation cadence and estate size (Compute RP: 250 reads / 5 min, 1,200 writes / hour per subscription). [Documented]
- **Phase 1 guardrails:** destructive VM-association transfers (Tier 3) are blocked; auto-decrease is operator-gated; forecast is advisory. [Decided]

- **Consumer-quota-under-shared-reservation (POC-001)** is the top unknown; readiness must validate consumer quota until proven. [Assumed]
-

3. Logical Component Architecture (Section 23)

ACRME is a set of cooperating services around a versioned state store and a policy (config-as-code) service. Components are deployable as modules of a single control application; boundaries are drawn for testability and independent scaling.

Component	Responsibility	Key requirements
Inventory Collector	Polls Azure (CRGs, reservations, quota/groupQuotas, usage, associations, zones) into freshness-stamped snapshots.	DAT-001..004, CAP-001/002, NFR-002
State Reconciler	Every cycle compares desired vs actual; enforces Allocated + Buffer floor; guarded scale-down; never deletes.	CAP-003/004/009/010, OPS-002
Placement & Scoring Engine	Validates exact region (or derives on exception); computes PS_Prod/PS_NonProd/PS_DR ; writes seed record.	PLC-001..010, REG-001..003
Quota-Pool Manager	Maintains the single governed pool; computes earmarks; allocates on demand; enforces Prod floor & DR earmark; exception two-group mode.	QUA-001..014, ADR-002
DR Orchestrator	Maintains SourceDestinationDRIndex ; sizes max-not-sum; runs standby activation waves & staged acquisition; fallback.	DR-001..019, ADR-005
Readiness API	Combined capacity+quota+freshness verdict as a machine-readable state to AEP.	RDY-001..004, INT-001..007
State Store	Versioned, optimistic-concurrency document store for entities, snapshots, operation/audit records.	DAT-001..006, NFR-007
Config / Scope-File Service	Versioned PlacementPolicy , region catalogue, scope files; decision-traceable.	GOV-003, DAT-005, REG-001
Forecast & Increase Service	Advisory Forecast_Quantity ; auto-increase triggers; oper-	CAP-005, ADR-004

Component	Responsibility	Key requirements
	ator-gated capacity/quota increase.	
Emergency Transfer Service	Tier 1/2/3 escalation during DR_EVENT_ACTIVE ; quota-neutral Tier 2; Tier 3 blocked in Phase 1.	ADR-003, Section 13 baseline
Observability Emitter	Metrics, alerts, dashboards incl. per-destination max-source coverage & source↔destination map.	OBS-001..005
Governance & Audit	RBAC, managed identity, approval gates, break-glass, immutable audit with before/after + correlation IDs.	GOV-001..009

T1 — Logical component architecture

```

flowchart TB
    subgraph External
        AEP[AEP / Provisioning]
        OPS[Operators / DR coordinator]
        AZ[Azure control planes:<br/>Compute / Reservations / groupQuotas]
        OBS[Observability platform]
    end
    subgraph ACRME
        API[Readiness API]
        PE[Placement & Scoring Engine]
        REC[State Reconciler]
        QPM[Quota-Pool Manager]
        DRO[DR Orchestrator]
        FIS[Forecast & Increase Service]
        ETS[Emergency Transfer Service]
        INV[Inventory Collector]
        CFG[Config / Scope-File Service]
        GOV[Governance & Audit]
        OE[Observability Emitter]
        STORE[(State Store<br/>versioned + optimistic concurrency)]
    end
    AEP -->|readiness query / mutations| API
    OPS -->|approvals, DR declaration| API
    API --> PE & QPM & DRO & REC
    PE --> STORE
    QPM --> STORE
    DRO --> STORE
    REC --> STORE
    FIS --> STORE
    ETS --> STORE
    INV -->|snapshots| STORE
    INV <--> AZ
    QPM <--> AZ
    REC <--> AZ
    DRO <--> AZ
    CFG --> PE & QPM & DRO & REC & FIS
    GOV -.enforces.-> API
    OE --> OBS
    STORE --> OE

```

4. Runtime & Deployment (Section 23, CAP-006, PLC-009)

4.1 Execution model

- **Reconciliation job (container app):** a long-running scheduler executes the reconciliation loop on a **6-minute target interval (configurable, CAP-006)**; targeted/critical reconciles run out-of-band. It drives the Inventory Collector → State Reconciler → Quota-Pool Manager cycle and emits snapshots. [Decided]
- **Placement flow (function/event path, PLC-009):** onboarding and AEP readiness requests are handled synchronously against the latest snapshot; on stale state the request triggers a synchronous refresh or returns `STALE_STATE`. [Decided]

- **DR path:** DR declaration flips `engine_mode` to `DR_EVENT_ACTIVE` ; the DR Orchestrator drives priority-wave activation out of the steady-state cadence. [Decided]

4.2 Cadence, SLOs and drift (Section 38, NFR)

```
Reconciliation loop target      = 6 min (configurable; production interval TBD)
[Assumed]
Critical targeted reconcile P95 < 2 min
[Assumed]
Stable-resource state age P95 < 15 min
[Assumed]
Drift detection                 ≤ 2 reconciliation cycles
[Assumed]
Snapshot max age (staleness gate) = max_snapshot_age_seconds (config, RDY-004)
[Decided]
```

4.3 API surface (summary; detail in Section 11)

A readiness/placement API (query + idempotent mutations), an operations/approval API (increase approval, DR declaration, override, sharing activation), and a read-only audit/observability API. All mutating operations are idempotent and versioned (INT-004..006). [Decided]

T2 — Management group / subscription topology

```
flowchart TB
    MG[Management Group root] --> MGPlat[Platform MG]
    MG --> MGWork[Workload MGs]
    MGPlat --> SubCtl[ACRME control subscription<br/>engine, state store, config]
    MGPlat --> SubObs[Observability subscription]
    MGWork --> SubProd[Prod subscriptions<br/>Prod CRGs + pool]
    MGWork --> SubNonProd[NonProd/CVAL subscriptions<br/>NonProd CRGs]
    MGWork --> SubDR[DR-hosting subscriptions<br/>DR standby CRGs]
    SubCtl -. managed identity + RBAC .-> SubProd
    SubCtl -. managed identity + RBAC .-> SubNonProd
    SubCtl -. managed identity + RBAC .-> SubDR
```

5. Topology — Subscriptions, CRGs and Quota Groups (Section 24, Section 25, Section 26)

5.1 Management group / subscription model (Section 24)

ACRME runs from a platform control subscription with cross-subscription managed-identity access (least-privilege RBAC) to the Prod, NonProd/CVAL, and DR-hosting subscriptions (see T2). All mutations are attributed to the control identity and audited (GOV-002). [Decided]

5.2 CRG hierarchy and cross-subscription sharing (Section 25, FC-06)

- **Environment-scoped, region + per-AZ CRG structure (CAP-023, extends CAP-011):** each environment (Prod, NonProd/CVAL, DR-standby) owns, per region, **one regional CRG plus one CRG per availability zone** in that region. The regional CRG anchors region-scoped (non-zonal-pinned) reservations and governance; the per-AZ CRGs hold the zonal `seed` matrix and zone-pinned reservations that back the even-distribution target (PLC-011). This refines the earlier “three-CRG-per-region” shorthand: the three remains the count of **environment** targets per re-

gion, each of which now expands into 1 regional + N per-AZ CRGs. Naming follows OPS-006 (e.g. `crg-pr-eus2-reg`, `crg-pr-eus2-az1`, `crg-pr-eus2-az2`, `crg-pr-eus2-az3`). [Decided]

- **Zone alignment (FC-06):** cross-subscription CRG sharing requires a validated `ZoneMappingRecord` — logical zones differ across subscriptions, so sharing is only valid when physical zones align (CAP-014..019). [Decided]
- **Sharing limits:** provider→consumer sharing validated for SKU/region/zone/authorisation; a per-provider consumer ceiling applies. [Decided]

T3 — CRG hierarchy & cross-subscription sharing (zone alignment)

```

flowchart TB
    subgraph RegionR [Region R]
        ProdCRG[Prod CRG]
        NonProdCRG[NonProd/CVAL CRG]
        DRCRG[DR-standby CRG]
    end
    ProdCRG -->|reserves| ZmapP["ZmapP{{Zone map: logical->physical}}"]
    NonProdCRG --> ZmapN["ZmapN{{Zone map}}"]
    DRCRG --> ZmapD["ZmapD{{Zone map}}"]
    ZmapP --> ShareA["ShareA[Consumer subscription A<br/>zone-aligned share]"]
    ZmapN --> ShareB["ShareB[Consumer subscription B]"]
    ZmapD --> ShareC["ShareC[DR consumer subscription]"]
    note1["FC-06: share valid only when physical zones align<br/>ZoneMappingRecord required"]:::n
    classDef n fill:#fff3cd,stroke:#e0a800;

```

5.3 Quota-group architecture — single-pool decision (Section 26, QUA-004)

Primary model: one governed quota group per region and quota family, covering Prod + NonProd/CVAL + DR together. All eligible default per-region quota is collected into the pool (QUA-003) and allocated on demand. Prod and DR are protected inside the shared pool by **logical earmarks** enforced by the Quota-Pool Manager at allocation/reclamation time — not by physical group separation. Because all three environments share one pool, an emergency DR draw needs **no cross-group transfer**. [Decided]

Exception topology: two (or more) groups per region, used **only** when Azure Quota Group limits or a mandatory Prod-isolation governance boundary make one pool impossible; each case is recorded in the Decision Log and reverts to one pool when the constraint lifts (DEP-001). [Decided]

The pool arithmetic (`Pool_Limit`, `Prod_Reserved_Floor`, `DR_Earmark_vCPU`, `Allocatable_NonProd`, `Pool_Headroom`) is specified in Section 8.3. See T4.

T4 — Quota-group / pool architecture (single-pool model)

```

flowchart TB
    subgraph Pool [Single governed quota pool per region/quota family]
        direction TB
        PF[Prod_Reserved_Floor<br/>= Prod_Used + Prod_Growth_Buffer<br/>never allocat-  
able to NonProd]
        DE[DR_Earmark_vCPU<br/>= Destination_DR_Requirement x vCPU max-not-sum<br/>  
>never allocatable to NonProd]
        AN[Allocatable_NonProd<br/>= Pool_Limit - Prod_Reserved_Floor - DR_Earmark -  
NonProd_Used]
        PH[Pool_Headroom = Pool_Limit - Pool_Used]
    end
    Pool -->|emergency DR draw: pool headroom first,<br/>then reclaimable NonProd - no  
cross-group transfer| DR[DR activation]
    Exc[Exception topology only:<br/>Prod group + NonProd+DR group<br/>when Azure lim-  
it / Prod-isolation boundary]:::e
    classDef e fill:#f8d7da,stroke:#dc3545;

```

6. Data Architecture (Section 34, DAT-001..006)

6.1 Store characteristics

A versioned document store with per-document optimistic concurrency (`policy_version` / `_etag` -style guards), freshness metadata on every snapshot, and append-only operation/audit records (DAT-001..006, NFR-007). Every decision-driving read carries a `capacity_snapshot_ref` for deterministic replay. [Decided]

6.2 Core entities

The three baseline-critical entities are `CustomerSeedRecord`, `SourceDestinationDRIndex`, and `CVALEarmarkRecord`, alongside reservation/quota snapshots, `PlacementPolicy`, and operation/audit records.

```

CustomerSeedRecord {                                # PLC-003/004/005 – first-placement author-
ity
    customer_realm_id      : string (PK)
    geography              : string (PK)
    production_region      : string
    cval_region            : string
    dr_region              : string | "NOT_OFFERED"
    products_covered       : list<string>
    decision_timestamp      : datetime
    policy_version         : string
    capacity_snapshot_ref  : string
    exception_ref          : string | null
    approval_metadata      : ApprovalRecord
    input_mode             : "SPECIFIC_REGION" | "GEOGRAPHY_EXCEPTION"
}

SourceDestinationDRIndex {                          # DR-018 – reverse-of-seed, drives activa-
tion
    source_region          : string (PK)
    destination_region     : string (PK)
    customer_realm_id      : string (PK)
    standby_instance_set   : list<InstanceRef>
    sku_family             : string
    quantity_vcpu          : int
    activation_state       : STANDBY | ACTIVATING | ACTIVE | FAILBACK_PENDING
    last_updated           : datetime
    policy_version         : string
}

CVALEarmarkRecord {                                # PLC-010 – prevents DR/CVAL double-count
    customer_realm_id      : string (PK)
    region                 : string (PK)
    cval_capacity_vcpu     : int
    dr_earmarked_vcpu      : int                    # portion counted toward DR, not live CVAL
    headroom
        co_located         : bool
        last_updated       : datetime
}

```

Supporting entities: `RegionalSnapshot` (freshness-stamped capacity/quota/usage/zone data), `PlacementPolicy` (weights, buffers, thresholds, catalogue — versioned), `ReservationState` / `QuotaPoolState`, `OperationRecord` (saga + compensation), `AuditEvent` (append-only), `ActivationRecord` (DR-019, Section 9.3), `ZoneMappingRecord` (FC-06), `SharingRelationship`, `DRD-distributionPlan`.

T5 — Data model / ERD (incl. seed record, SourceDestinationDRIndex, CVAL earmark)

```

erDiagram
    CustomerSeedRecord ||--o{ SourceDestinationDRIndex : "reverse view"
    CustomerSeedRecord ||--o{ CVALEarmarkRecord : "co-location"
    CustomerSeedRecord }o--|| PlacementPolicy : "policy_version"
    RegionalSnapshot ||--o{ QuotaPoolState : "captures"
    QuotaPoolState ||--o{ ReservationState : "governs"
    SourceDestinationDRIndex ||--o{ ActivationRecord : "activates"
    DRDistributionPlan ||--o{ SourceDestinationDRIndex : "aggregates"
    OperationRecord ||--o{ AuditEvent : "emits"
    CustomerSeedRecord {
        string customer_realm_id PK
        string geography PK
        string production_region
        string cval_region
        string dr_region
        string input_mode
        string policy_version
    }
    SourceDestinationDRIndex {
        string source_region PK
        string destination_region PK
        string customer_realm_id PK
        int quantity_vcpu
        string activation_state
    }
    CVALEarmarkRecord {
        string customer_realm_id PK
        string region PK
        int cval_capacity_vcpu
        int dr_earmarked_vcpu
        bool co_located
    }
    QuotaPoolState {
        string region PK
        int pool_limit
        int prod_reserved_floor
        int dr_earmark_vcpu
        int allocatable_nonprod
    }

```

6.3 Freshness & versioning

Every snapshot carries `collected_at`; the readiness gate compares `now() - collected_at` against `max_snapshot_age` (RDY-004). Policy and scope files are versioned and decision-traceable; the version that drove any decision is persisted with it (DAT-005, GOV-003). [Decided]

7. State Model & Concurrency (Section 29, NFR-002/007)

7.1 Engine mode machine (five states)

```
STEADY_STATE → DR_DECLARATION_PENDING → DR_EVENT_ACTIVE → FAILBACK_PENDING →
STEADY_STATE
                                     ↘ INCIDENT_HOLD (safe degraded)
```

Steady-state reconciliation and auto-increase run **only** in `STEADY_STATE` ; emergency transfer is callable **only** in `DR_EVENT_ACTIVE` . Entering `DR_EVENT_ACTIVE` does not auto-authorise service-impacting CVAL action — each action retains its own policy gate. `INCIDENT_HOLD` is the safe degraded state entered when inputs cannot be trusted. [Decided]

7.2 Concurrency & reservation-of-intent

- **Optimistic concurrency:** every mutating operation reads an expected version and fails the write on mismatch, forcing a re-read (NFR-007). [Decided]
- **Reservation-of-intent:** placement writes an atomic hold before the seed record so two concurrent onboardings cannot double-commit the same capacity (NFR-002). [Decided]
- **Saga + compensation:** multi-step mutations use `OperationRecord` with a compensation chain for rollback; `VM_ImpactRecord` is immutable for any VM state change. [Decided]

T6 — Five-state engine machine

```
stateDiagram-v2
    [*] --> STEADY_STATE
    STEADY_STATE --> DR_DECLARATION_PENDING: DR declared (authorised)
    DR_DECLARATION_PENDING --> DR_EVENT_ACTIVE: activation approved
    DR_DECLARATION_PENDING --> STEADY_STATE: declaration withdrawn
    DR_EVENT_ACTIVE --> FAILBACK_PENDING: failback authorised
    FAILBACK_PENDING --> STEADY_STATE: failback complete (reversible)
    STEADY_STATE --> INCIDENT_HOLD: inputs untrusted / stale
    DR_EVENT_ACTIVE --> INCIDENT_HOLD: safety trip
    INCIDENT_HOLD --> STEADY_STATE: inputs restored
```

8. Placement Scoring & Forecasting (Section 28)

8.1 Input modes and pipeline

The default input is an **exact production region** which ACRME **validates** (it does not derive). A **geography** input is an **exception path** requiring approval + customer acknowledgement, after which the derived region becomes the fixed seed (ADR-001, PLC-001/002). Restricted regions never enter scoring — they route to the exception workflow.

Reservation-eligibility gate (CAP-020/CAP-021). Before a VM enters zonal scoring the Onboarding Validator applies two hard preconditions: (1) **Availability-Set VMs are ineligible** for zonal capacity reservations (**CAP-020**) — a VM in an Availability Set cannot be backed by a zonal on-demand capacity reservation and is rejected from the zonal path; (2) a VM that is not already **zone-pinned** must satisfy the **deallocate/redeploy-to-AZ precondition (CAP-021)** — it is deallocated and redeployed into a target availability zone (via the governed onboarding/migration workflow) before it can be asso-

ciated with a per-AZ CRG reservation. Only zone-pinned, non-Availability-Set VMs proceed to the PLC-011 even-distribution zone selection (Section 8.5).

T7 — Staged placement pipeline

```

flowchart LR
    In[Placement request] --> HC[Hard constraints HC-1..HC-10]
    HC -->|fail| Rej[Reason state / exception]
    HC -->|pass| Sel{Input mode}
    Sel -->|exact region default| Val[Validate & accept]
    Sel -->|geography exception| Score[Score Standard regions PS_Prod]
    Score --> Pick[argmax -> fixed seed]
    Val --> Seq[Sequential CVAL then DR selection]
    Pick --> Seq
    Seq --> Seed[Write seed + DR index]

```

T8 — Prod region input modes (exact vs geography exception)

```

flowchart TD
    R[Region input] --> Q{Exact region or geography?}
    Q -->|Exact Standard| V[Validate HC-1..HC-10 -> accept, no argmax]
    Q -->|Exact Restricted| X[Exception Deployment Workflow]
    Q -->|Geography| E{Exception approved + acknowledged?}
    E -->|No| B[POLICY_BLOCKED]
    E -->|Yes| D[argmax PS_Prod over Standard regions -> fixed seed]

```

8.2 Scoring formulas (weights $\alpha=0.30/\beta=0.20/\gamma=0.25/\delta=0.15/\epsilon=0.10$)

All five weights are tunable `PlacementPolicy` constants summing to 1.0; every component is clamped `Clamp(x)=max(0,min(1,x))` before weighting. [Assumed]

PlacementPolicy weight-sum validation gate [Decided]

The weight-sum invariant ($\alpha + \beta + \gamma + \delta + \epsilon = 1.0$) is a **mathematical requirement**, not a convention. Without it, placement scores fall outside $[0, 1]$, rendering threshold gates and cross-version comparisons invalid regardless of which region `argmax` selects.

Responsibility: Config / Scope-File Service. This service is the sole gatekeeper for `PlacementPolicy` versions. The validation gate runs at **policy version load time**, before the version is published to the Placement & Scoring Engine. The Placement Engine never receives a policy version that has not passed this gate.

Formal gate logic:

```

weight_sum = policy.alpha + policy.beta + policy.gamma + policy.delta + policy.epsilon
tolerance = 0.001 # IEEE 754 floating-point rounding only; not a margin for operation error

IF |weight_sum - 1.0| > tolerance:
    REJECT version – do NOT publish
    RAISE PolicyValidationError(code="WEIGHT_SUM_INVARIANT_VIOLATED", computed_sum=weight_sum)
    WRITE AuditEvent(event_type="POLICY_VALIDATION_FAILED", severity="ERROR", policy_version=...)
ELSE:
    PUBLISH version to Placement Engine
    WRITE AuditEvent(event_type="POLICY_VERSION_PUBLISHED", weight_sum=weight_sum)

```

Consequences of a non-1.0 sum:

Weights sum to	Scoring effect	Operational consequence
> 1.0 (e.g. 1.10)	Best possible score > 1.0 — scores inflated	Placement thresholds pass regions that should be blocked
< 1.0 (e.g. 0.90)	Best possible score < 1.0 — scores compressed	Placement thresholds block regions that should pass; cross-version comparisons invalid

Tuning constraint: when any weight is increased, the remaining weights must decrease by the same total. The tuning produces a valid new sum before the updated policy version is submitted:

```

Valid:  α=0.40, β=0.15, γ=0.25, δ=0.10, ε=0.10 → sum = 1.00 ✓
Invalid: α=0.40, β=0.20, γ=0.25, δ=0.15, ε=0.10 → sum = 1.10 ✗ – gate rejects

```

Test catalogue entries (T-WV-01..T-WV-05): see Calculation Logic Reference Section “Policy weight-sum validation rule” for the full test matrix covering the valid case, both invalid-sum directions, IEEE 754 rounding tolerance, and the tolerance boundary.

[Decided]

```

PS_Prod(r)    = 0.30·Clamp(nonprod.effective_free / prod.quantity)
               + 0.20·Clamp(prod.quota_headroom / prod.quota_limit)
               + 0.25·Clamp(1 - prod_customer_count / total_customers(g))
               + 0.15·Clamp(dr.coverage_ratio)
               + 0.10·Clamp(az_count / 3)

PS_NonProd(r) = 0.30·Clamp(nonprod.effective_free / nonprod.quantity)
               + 0.20·Clamp(nonprod.quota_headroom / nonprod.quota_limit)
               + 0.25·Clamp(1 - nonprod_customer_count / total_customers(g))
               + 0.15·Clamp(nonprod.effective_free / nonprod.quantity)
               + 0.10·Clamp(az_count / 3)

PS_DR(r)      = 0.30·Clamp(dr.free_slots / dr.quantity)
               + 0.20·Clamp(dr.quota_headroom / dr.quota_limit)
               + 0.25·Clamp(1 - dr_customer_count / total_customers(g))
               + 0.15·min(1.0, dr.coverage_ratio / dr_ratio_target)
               + 0.10·Clamp(az_count / 3)

```

y denominator (PLC-012, A.10) [Decided] : `total_customers(g)` is the count of distinct customers with any environment (Prod, CVAL or DR) provisioned or being deployed in any region of the scored region's geography g. It is one live count per geography, shared by all three scores. **y = 1 when it is zero.**

Note on the PS_DR δ term: `dr.coverage_ratio` is `dr.quantity / Destination_DR_Requirement(region)`, where `Destination_DR_Requirement(region) = MAX(source portions)` (max-not-sum, A.6/DR-017). The `dr_ratio_target` divisor here is only a **configurable normalization reference for the coverage-health signal** — it is **not** the retired fixed 30–40% DR-to-Prod sizing ratio (Calc Logic Reference Scenario 15 → 17). DR sizing itself uses max-not-sum, never a fixed ratio.

`PS_Prod` is dual-purpose (derive on exception; validate/audit otherwise). Every candidate score + policy version + snapshot ref is written to the `OperationRecord` for deterministic replay. A reviewer-recommended pilot `PS_NonProd` variant (removing the α/δ duplication) is recorded in the Calculation Logic Reference; the formula above is the approved design-of-record. Scoring runs in shadow/recommendation mode until empirically validated. [Decided]

8.3 Quota-pool arithmetic (single-pool, Section 26/QUA-004)

```

Pool_Limit(region) = Prod_CRG_qty·vCPU·(1+prod_growth_buffer)
                   + NonProd_CRG_qty·vCPU·(1+nonprod_growth_buffer)
                   + DR_Earmark_vCPU(region)
                   + emergency_transfer_headroom_vcpu

Prod_Reserved_Floor = Prod_Used_vCPU + Prod_Growth_Buffer_vCPU           # never alloc-
atable to NonProd
DR_Earmark_vCPU      = Destination_DR_Requirement(region)·vCPU           # max-not-sum,
never allocatable
Allocatable_NonProd  = Pool_Limit - Prod_Reserved_Floor - DR_Earmark_vCPU - Non-
Prod_Used_vCPU
Pool_Headroom        = Pool_Limit - Pool_Used
Emergency_DR_Available = Pool_Headroom + reclaimable_NonProd_above_committed

```

NonProd allocation fail-safes when `Allocatable_NonProd` reaches zero. A dual-validation detector re-computes the DR earmark from authoritative assignment data; disagreement disables automatic NonProd expansion until reconciled (Scenario 9). [Decided]

8.4 Forecasting vs reconciliation floor

```

Target Reserved Capacity = Allocated VM Count + Configured Buffer      # CAP-003
continuous floor
Forecast_Quantity          = ceil(Forecast_Peak · (1+Growth_Buffer) + DR_Buffer) #
ADR-004 proactive growth

```

The reconciliation floor is enforced every cycle; the forecast is advisory (horizon $\in \{30,60,90\}$ days) and raises `ForecastApproachingQuotaLimit` with 14-day lead time at 80% of quota limit. Auto-increase target = max(floor, forecast). [Decided]

8.5 Even per-zone distribution & rebalancing (PLC-011, Baseline App. A.9)

Within a placement region with `Z` availability zones, ACRME drives the per-zone VM distribution toward an even target and raises a rebalancing action when it drifts:

```

zone_target_share      = 1 / zone_count      # e.g. 0.333 for
Z=3 (C-13)
zone_share(z)          = vm_count(z) / total_vm_count
zone_deficit(z)        = zone_target_share - zone_share(z)      # >0 means under-
target
# Placement zone selection (per request):
eligible(z)            = has_capacity(z) AND has_quota(z) AND NOT restricted(z) AND
zone_aligned(z)
chosen_zone            = argmax_{z in eligible} zone_deficit(z)  # greatest-defi-
cit-first
# Drift detection (per reconciliation cycle):
max_drift              = max_z | zone_share(z) - zone_target_share |
rebalance_needed       = max_drift > zone_balance_tolerance      # C-13 default
0.10 (±10 pp)

```

When `rebalance_needed`, a `ZoneRebalanceAction` is raised so subsequent placements (and, where policy permits, governed redeploys) steer new/moved VMs into the most-deficient eligible zones until `max_drift ≤ tolerance`. Rebalancing never violates capacity, quota, restriction, or zone-alignment constraints and never forces a live migration outside an approved workflow. Both `zone_target_share` (derived from `zone_count`) and `zone_balance_tolerance` are configuration-driven (**C-13**); the worked example is Baseline Appendix A.9. [Decided]

9. DR Sizing & Activation Algorithms (Section 31, App. A.6/A.7/A.8, App. D)

9.1 Max-not-sum sizing (DR-017, ADR-005)

```

Destination_DR_Requirement(d) = MAX over each non-concurrent source s protected by d
                               ( Workload_Portion(s -> d) )
DR_Capacity_Gap(d)            = max(0, Destination_DR_Requirement(d) - Us-
able_Destination_Capacity(d))
Overcommit_Ratio(d)           = SUM(Workload_Portion(s->d)) / MAX(Workload_Portion(s-
>d))

```

Single-failure assumption (DR-001): one production region fails at a time, so destination `d` needs standby for its **largest** protected source only; standby is shared/overcommitted across mutually exclusive events. `Usable_Destination_Capacity(d)` may include approved bootstrap headroom, available CRG reservations, releasable CVAL (after earmark check), and approved sharing/expansion. A per-scope **SUM override (C-11)** covers contractual simultaneous-failure protection with no code change.

[Decided]

[Amended v2.4] Middle East sources contribute a cross-geo DR term at their Europe destination (DR-020, DEC-001 RESOLVED). Middle East sources are now DR-protected cross-geo into a weighted-selected Europe Standard region. Each Middle East source therefore contributes a `Workload_Portion(s → d)` term at its Europe destination, produces a `SourceDestinationDRIndex` entry, and reserves `DR_Earmark_vCPU` **at that Europe destination** — where the **max-not-sum** rule (DR-017) applies across all sources (Middle East and any European sources) that fail over to the same Europe region. This adds DR load to Europe; Europe destinations must size for the maximum concurrent DR requirement, not the sum. Any remaining `DR_NOT_OFFERED` geography (a future config state, not the Middle East) would still be excluded from sizing as before. [Documented]

T12 — Max-not-sum sizing (overcommit ratio)

```
flowchart TB
    S1[Source R1 -> R2: 120 cores]
    S3[Source R3 -> R2: 80 cores]
    S1 --> R2
    S3 --> R2
    R2[Destination R2 standby<br/>SUM = 200 cores over-provisions<br/>MAX = 120 cores sized<br/>saving = 80 cores 40%<br/>Overcommit_Ratio = 200/120 ≈ 1.67]:::d
    classDef d fill:#d4edda,stroke:#28a745;
```

9.2 Standby activation waves & staged acquisition (DR-006, DR-019)

On authorised declaration the DR Orchestrator reads the `SourceDestinationDRIndex` for the failed source and activates customers in business-priority waves (P0 platform → P-1 highest → P-N), transitioning each `associated → allocated`. Per wave, staged acquisition runs in order:

```
Stage 1: approved bootstrap / pre-staged headroom (zero Azure wait)
Stage 2: available reservation + quota already in destination
Stage 3: shut down / disassociate eligible CVAL (after DR-010 authorised trigger)
Stage 4: share/reassign reservations within supported region/zone boundaries
Stage 5: allocate pooled quota to DR subscriptions
Stage 6: request additional Azure quota/capacity where required
Stage 7: report unrecoverable capacity gaps
```

Stages 1-2 start immediately from the single pool's headroom, so P0/P-1 recovery does not wait on throttled Azure deallocation APIs. [Decided]

9.3 Activation state tracking & failback (DR-013)

```

ActivationRecord {
  customer_realm_id, source_region, destination_region,
  priority_wave : int,
  activation_state : PENDING | ACTIVATING | ACTIVE | FAILED | FAILBACK_PENDING,
  stage_reached : int (1..7), capacity_gap_vcpu : int,
  activated_at, approved_by, incident_id
}

```

Failback reverses `ACTIVE` → `FAILBACK_PENDING` → `STANDBY` in reverse priority order; seed record and DR index are preserved (no re-derivation). `[Decided]`

T11 — DR standby activation sequence (index lookup → waves → staged acquisition)

```

sequenceDiagram
    participant DRC as DR coordinator
    participant DRO as DR Orchestrator
    participant IDX as SourceDestinationDRIndex
    participant QPM as Quota-Pool Manager
    participant AZ as Azure
    DRC->>DRO: Declare source-region failure (authorised)
    DRO->>DRO: engine_mode = DR_EVENT_ACTIVE
    DRO->>IDX: Query standby sets for failed source
    IDX->>DRO: Customers + priority waves
    loop each priority wave (P0 -> P-1 -> P-N)
        DRO->>QPM: Stage1-2 draw pool headroom + reservation
        QPM-->>DRO: Earmark released (no cross-group transfer)
        alt deficit remains
            DRO->>QPM: Stage3 CVAL release / Stage4 share / Stage5 pooled quota
            DRO->>AZ: Stage6 request additional capacity/quota
            AZ-->>DRO: Granted or deficit
        end
        DRO->>DRO: associated -> allocated (ActivationRecord)
        DRO->>DRC: Wave complete + audit checkpoint
    end
    DRC->>DRO: Failback authorised
    DRO->>DRO: ACTIVE -> FAILBACK_PENDING -> STANDBY (reverse order)

```

9.4 Tier escalation (Section 32) — see Section 13 baseline

Emergency capacity during `DR_EVENT_ACTIVE` escalates Tier 1 (direct expansion) → Tier 2 (quota-neutral transfer within the shared pool) → Tier 3 (destructive, dual-approval, **blocked in Phase 1**). In the single-pool model Tier 2 is intrinsically quota-neutral (draw and release hit the same pool). `[Decided]`

T10 — Three-tier emergency transfer escalation

```

flowchart TD
    Ev[DR event active] --> T1{Tier 1 headroom available?}
    T1 -->|Yes| A1[DirectExpansion - automated additive]
    T1 -->|No| T2{Tier 2 approved?}
    T2 -->|Yes| A2[QuotaNeutralTransfer - reduce NonProd, expand DR in same pool]
    T2 -->|No| T3{Tier 3 allowed?}
    T3 -->|Phase 1: BLOCKED| M[Manual escalation]
    T3 -->|Later phase: dual-approval| A3[DestructiveTransfer - changes VM associations]

```

T13 — Distributed DR reference model (Section 12A)

Reused reference diagram (ADR-005 Section 12A): `adr/diagrams/ac-rme_three_region_capacity_model.png` — each region concurrently hosts Prod + CVAL + DR standby tagged with its `src Rn` source; each destination is sized by max-not-sum against its largest single protected source.

10. Steady-State Capacity Lifecycle (Section 30)

The normative 10-step sequence runs only in `STEADY_STATE` :

1. Detect threshold crossing (utilisation $\geq$ auto-increase threshold).
2. Re-read current CR, quota, sharing, assignment state.
3. Create CapacityIncreaseRequest.
4. Calculate target = max(Allocated+Buffer floor, Forecast_Quantity).
5. Require operator approval (Phase 1).
6. Submit quota action only if validated as required.
7. Wait for confirmed quota state — no assumed propagation SLA.
8. Update CR quantity.
9. Confirm actual quantity.
10. Refresh snapshot and close the request.

Triggers: `dr_autoincrease_threshold=0.35` , `prod/nonprod_autoincrease_threshold=0.20` , debounce cooldown 30 min per (region+CRG type). Guarded scale-down never drops below `Allocated` , within a DR window, during maintenance exclusion, or against cost policy. **Decommissioning-workflow boundary (CAP-008/CAP-010):** automatic reconciliation only ever **right-sizes a reservation down to its `Allocated + Buffer floor`** — it never reduces a reservation to zero and never deletes a CRG or reservation on its own. Reducing a reservation to **zero (returning it to a `seed`)**, **retiring**, or **deleting** a reservation/CRG are **gated workflow actions** requiring the decommissioning approval path (operator approval, cost/DR/maintenance guards, audit) — they are outside the automatic loop. [Decided]

T9 — Steady-state 10-step capacity lifecycle

```
sequenceDiagram
    participant REC as State Reconciler
    participant OP as Operator
    participant QPM as Quota-Pool Manager
    participant AZ as Azure
    REC->>REC: 1. Detect threshold crossing
    REC->>AZ: 2. Re-read CR/quota/sharing/assignment
    REC->>REC: 3. Create CapacityIncreaseRequest
    REC->>REC: 4. Target = max(Allocated+Buffer, Forecast)
    REC->>OP: 5. Require approval (Phase 1)
    OP-->>REC: Approved
    REC->>QPM: 6. Submit quota action if required
    QPM->>AZ: 7. Await confirmed quota state
    AZ-->>QPM: Confirmed
    QPM->>AZ: 8. Update CR quantity
    AZ-->>QPM: 9. Confirm actual quantity
    REC->>REC: 10. Refresh snapshot + close request
```

11. API Architecture (Section 35, INT-001..007)

11.1 Contracts

```

POST /readiness          → DeploymentReadinessResult  (query; idempotent)
  in : { customer_realm_id, region (exact), sku_family, environment, product, inten-
  ded_demand }
  out: { readiness_state, reasons[], operation_handle, snapshot_ref, policy_version }

POST /placement          → PlacementResult            (idempotent; writes seed + DR
index)
POST /capacity/increase  → OperationRecord          (operator-gated; idempotent)
POST /dr/declare         → OperationRecord          (authorised; flips en-
gine_mode)
POST /dr/activate-wave   → ActivationRecord[]         (per priority wave)
POST /quota/allocate     → OperationRecord          (single-pool allocation)
POST /sharing/activate   → OperationRecord          (zone-aligned share)
GET  /audit, GET /state   → read-only

```

11.2 Idempotency, versioning, error/readiness codes

- Every mutating call carries an **idempotency key** and an **expected-state version**; retried calls re-
turn the original result, mismatched versions return a concurrency error (INT-004..006, NFR-007).
[Decided]
- Readiness states are the machine-readable contract to AEP:

```

READY | READY_WITH_RISK | QUOTA_DEFICIT | RESERVATION_DEFICIT |
CAPACITY_UNAVAILABLE | STALE_STATE | POLICY_BLOCKED | VALIDATION_REQUIRED

```

- AEP **polls** operation handles rather than assuming synchronous completion, and never treats
provider quota as proof of consumer deployability (INT-007, POC-001). [Assumed]

11.3 Readiness gate logic (RDY-001)

```

READY requires ALL of:
1. target_region ∈ approved_region_catalogue          (REG-001)
2. az_count ≥ 1 in target_region                      (CAP-011)
3. sku ∈ supported_skus(target_region)
4. reservation_policy known for (subscription, region, zone, sku)
5. reservation exists (if required by policy)
6. reserved_quantity ≥ requested_count OR approved_over_allocation active
7. consumer_subscription_quota ≥ requested_units      (QUA-013 – POC-gated)
8. snapshot_age ≤ max_snapshot_age_seconds           (RDY-004)
9. no blocking policy exception or hard-constraint violation
Quota_Deficit = max(0, Requested_Deployment_Units - Available_Quota)
Available_Quota = Assigned_Regional_VM_Family_Quota - Current_Regional_VM_Family_Usage

```

12. Security (Section 36, GOV-001..009)

- **RBAC & managed identity:** the control identity holds least-privilege, scoped roles per target
subscription; no standing user credentials for mutations (GOV-001/002). [Decided]

- **Approval gates:** capacity/quota increase, DR declaration, CVAL release, Tier 2 transfer, seed change, and overrides are policy-gated with recorded approver identity (GOV-005..007). [Decided]
- **Break-glass:** an audited emergency-elevation path with dual control and time-boxed elevated RBAC; every break-glass action is immutably logged (GOV-008). [Decided]
- **Secrets:** managed identity + platform secret store; no secrets in config or code. [Decided]
- **Audit:** append-only `AuditEvent` for every decision and mutation with before/after values and correlation IDs (GOV-004/009, DAT-005). [Decided]

13. Observability Implementation (Section 37, OBS-001..005)

Signal	Metric / alert	Requirement
Buffer deficit	<code>reservation_buffer_deficit</code> gauge; alert when reserved < <code>Allocated+Buffer</code>	OBS-001, CAP-003
Per-destination DR coverage	<code>dr_max_source_coverage(d)</code> ; alert when usable < <code>Destination_DR_Requirement(d)</code> (the MAX, not the SUM)	OBS-002, DR-017
DR capacity gap	<code>dr_capacity_gap(d)</code> gauge	OBS-002
Source↔destination mapping	dashboard of the <code>SourceDestinationDRIndex</code>	OBS-004
Overcommit ratio	<code>overcommit_ratio(d)</code> ; alert above safety ceiling	FIN-006..008, POC-011
Activation state	per-wave <code>ActivationRecord</code> progress	OBS-003, DR-019
Idle cost	idle reservation cost surfaced	OBS-005, FIN-001/002
Staleness	<code>snapshot_age</code> ; alert/refresh on breach	RDY-004, NFR-002
Forecast lead-time	<code>ForecastApproachingQuotaLimit</code> (14-day lead at 80%)	OBS-001, ADR-004

All metrics carry region/environment/scope dimensions and policy version for replay. [Decided]

14. NFRs & Resilience (Section 38, NFR-001..010)

- **Throttling resilience:** adaptive throttle manager initialises with ARM baselines (250 reads/5 min, 1,200 writes/hour per subscription) and adapts to observed 429/Retry-After ; reconciliation cadence and estate size are bounded accordingly (Section 14 calc). [Documented]
 - **Estate sizing:** $CRG_total = R \cdot Eclass \cdot Z \cdot SKUset \cdot IsolationFactor$; Requests_per_cycle summed across read classes / 300 s cycle to stay within budget. [Derived]
 - **Degraded mode:** on untrusted/stale inputs the engine enters INCIDENT_HOLD / returns STALE_STATE and fails safe rather than driving placement from stale data (NFR-002). [Decided]
 - **Simulation / dry-run:** placement, reconciliation, and **DR failover** can run in dry-run/simulation, producing the would-be decisions and readiness states without mutating Azure (NFR-006). [Decided]
 - **Idempotency & concurrency:** all mutations idempotent + optimistic concurrency (NFR-007). Determinism/replay from OperationRecord . [Decided]
-

15. Integration (Section 33, Section 6 review)

- **AKS / VMSS:** reservations back node pools / scale sets; ACRME readiness gates their capacity before scale-out; associations are tracked for reconciliation. [Decided]
 - **AEP:** the sole provisioning consumer of readiness; interacts only through the idempotent readiness/placement contract (Section 11) and polls operation handles. [Decided]
 - **Azure control planes:** Compute (CRGs/reservations), Reservations, and **groupQuotas** (single-pool). Single-pool enforcement depends on groupQuotas / groupType maturity (DEP-001). [Documented]
-

16. POC & Validation Dependencies (Section 42)

POC	Question	Gates	Status
POC-001	Does provider quota under a shared reservation guarantee consumer deployability?	Readiness gate step 7; INT-007	Top unknown
POC-006	Distributed DR topology behaves as modelled	ADR-005 Section 12A; DR-016	Required
POC-007	Lean bootstrap sizing is sufficient to initiate recovery	DR-007 bootstrap targets	Required
POC-011	max-not-sum over-commit is safe at platform scale	DR-017 sizing; over-commit safety ceiling; FIN-006..008	Required before production dependency
POC-005	VM state semantics (associated/allocated/deallocated)	DR-019 activation staging	Dependency
POC-001 (Tier 2 use)	Tier 2 quota-neutral transfer within the single pool: after a Quota Group transfer, does the deploying (consumer) subscription hold enough VM-family and total regional vCPU? (QUA-003, QUA-004, QUA-007)	ADR-003 “Relationship to Emergency Tiers”; Section 9 Tier 2	Required — Tier 2 stays approval-gated until proven (ADR-003)

Production reliance on max-not-sum DR sizing and single-pool consumer-quota behaviour is **gated** on POC-011 and POC-001 respectively; until then, readiness validates consumer quota explicitly and DR sizing carries the SUM-override safety option. [Decided]

17. Traceability — Requirement/Deviation → Component/Algorithm/Entity

Requirement group (v2.4 IDs)	Component(s)	Algorithm / formula	Entity / diagram
REG-001..003	Config/Scope-File, Placement Engine	catalogue validation; per-geography <code>distribution_model</code> ; input modes	PlacementPolicy ; T8
PLC-010a	Placement Engine, DR Orchestrator	two-region CVAL/DR co-location enforcement (EU/Australia/Asia Pacific)	CVALEarmarkRecord ; T7
PLC-010b [Amended v2.4]	Placement Engine, DR Orchestrator	Middle East cross-geo override: Prod+CVAL co-locate in a weighted-selected Middle East region, DR placed cross-geo in a weighted-selected Europe region (overrides PLC-010a for the Middle East)	CVALEarmarkRecord ; SourceDestination-DRIndex ; T7
ENV-001..007	Placement Engine, DR Orchestrator	env separation; role flip	CustomerSeedRecord ; T7
CAP-001, CAP-001a	Inventory Collector, Config/Scope-File	environment classification; core subscription = all-production (CAP-001a)	ReservationState , PlacementPolicy ; T2
CAP-001..019	Inventory Collector, State Reconciler	Allocated+Buffer floor; auto right-size only (no auto zero/delete); over-alloc	ReservationState ; T9, F5
CAP-008, CAP-010	State Reconciler, Decommissioning Workflow	auto right-size to <code>Allocated+Buffer</code> vs gated reduce-to-zero/retire/delete	ReservationState , OperationRecord ; Section 10, T9
CAP-020, CAP-021	Placement Engine, Onboarding Validator	Availability-Set VM ineligibility; deallocate/redeploy-to-AZ precondition	PlacementPolicy , CustomerSeedRecord ; Section 8.1
		seed matrix of count-0 reservations	

Requirement group (v2.4 IDs)	Component(s)	Algorithm / formula	Entity / diagram
CAP-022 (extends CAP-009)	Inventory Collector, Config/Scope-File	over eligible SKU×region×AZ; per-product-team budget governance	ReservationState (seed); Section 5.2
CAP-023 (extends CAP-011)	Inventory Collector, State Reconciler	regional + per-AZ CRG structure per environment	ReservationState (1 regional + N per-AZ CRGs); Section 5.2, T3
CAP-024 (reconciles CAP-019)	Inventory Collector, Governance & Audit	reactive SKU/AZ discovery auto-creates reservation + raises scope-file governance item	ReservationState , AuditEvent ; Section 5.2, Section 12
PLC-011	Placement Engine, State Reconciler	even ≈1/ zone_count target; greatest-deficit-first zone pick; drift → ZoneRebalanceAction	PlacementPolicy ; Section 8.5
OPS-006	Config/Scope-File, Provisioning	deterministic RG/ CRG/subscription naming convention + counter	PlacementPolicy ; Section 5.1, Section 5.2
C-12, C-13	Config Service	naming pattern config (C-12); zone target/tolerance config (C-13)	config; Section 8.5
QUA-001..014	Quota-Pool Manager	Section 8.3 pool arithmetic; earmarks; quota-as-governor	QuotaPoolState ; T4
RDY-001..004	Readiness API	Section 11.3 gate logic; staleness; quota deficit	readiness states; T6
PLC-001..010	Placement & Scoring Engine	PS_Prod/NonProd/DR ; seed reuse; co-location guard	CustomerSeedRecord , CVALEarmarkRecord ; T7, T8
DR-001..019	DR Orchestrator, Emergency Transfer	max-not-sum; gap; overcommit; activa-	SourceDestination-DRIndex ,

Requirement group (v2.4 IDs)	Component(s)	Algorithm / formula	Entity / diagram
		tion waves; staged acquisition; tiers	ActivationRecord ; T11, T12, T13, T10
FIN-001..008	Observability Emitter, Forecast Service	idle cost; overcommit ratio; cost-before-expansion	metrics; T12
INT-001..007	Readiness API	idempotency keys; expected-version; polling	OperationRecord ; T1
DAT-001..006	State Store, Config Service	freshness metadata; versioning; audit	all entities; T5
OBS-001..005	Observability Emitter	max-source coverage; source↔dest map; activation progress	dashboards; Section 13
GOV-001..009	Governance & Audit	RBAC; approval gates; break-glass; immutable audit	AuditEvent ; Section 12
NFR-001..010	(cross-cutting)	throttle resilience; degraded mode; idempotency; simulation	OperationRecord ; Section 14
OPS-001..005	(runbooks)	DR declaration/activation/failback; quota allocation; sharing; override	ActivationRecord ; Section 10, Section 16
POC-001..011	(gating)	consumer quota; DR topology; bootstrap; overcommit safety	Section 16
DEC-001..003 / DEP-001	Config/Scope-File, Quota-Pool Manager	Middle East DR policy — [Amended v2.4] DEC-001 RESOLVED: Middle East DR is offered cross-geo into a weighted-selected Europe Standard region (Prod+CVL	PlacementPolicy ; Section 2.2, Section 5.3

Requirement group (v2.4 IDs)	Component(s)	Algorithm / formula	Entity / diagram
		co-located in a Middle East region, DR reserved in Europe, weighted model, DR-020); failback duration; geo-exception approver; groupQuotas maturity	

Every Baseline v2.4 requirement group — including the reservation-model additions CAP-001a, CAP-020..024, PLC-011, OPS-006 and configuration items C-12/C-13 — resolves to at least one component, algorithm, and entity above. Functional-level traceability is in the FDD Section 9.

18. Class Diagram (updated with new entities)

T14 — Domain class model (updated: +CustomerSeedRecord, +SourceDestinationDRIndex, +CVALEarmarkRecord; single-pool QuotaPoolState ; readiness enum)

```

classDiagram
class CustomerSeedRecord {
    +string customer_realm_id
    +string geography
    +string production_region
    +string cval_region
    +string dr_region
    +string input_mode
    +string policy_version
}
class SourceDestinationDRIndex {
    +string source_region
    +string destination_region
    +string customer_realm_id
    +int quantity_vcpu
    +ActivationState activation_state
}
class CVALEarmarkRecord {
    +int cval_capacity_vcpu
    +int dr_earmarked_vcpu
    +bool co_located
}
class QuotaPoolState {
    +int pool_limit
    +int prod_reserved_floor
    +int dr_earmark_vcpu
    +int allocatable_nonprod
    +int pool_headroom
}
class PlacementPolicy {
    +float alpha_beta_gamma_delta_epsilon
    +float prod_growth_buffer
    +float nonprod_growth_buffer
    +RegionCatalogue region_catalogue
    +string version
}
class RegionCatalogue {
    +string geography
    +string distribution_model
    +list~string~ regions
    +list~string~ restricted_regions
    +bool dr_not_offered
}
class ActivationRecord {
    +int priority_wave
    +int stage_reached
    +int capacity_gap_vcpu
    +ActivationState state
}
class ReservationState {
    +int reserved_quantity
    +int allocated_count
}
class OperationRecord {
    +string operation_handle
    +string expected_version
}
CustomerSeedRecord "1" --> "0..*" SourceDestinationDRIndex : reverse view
CustomerSeedRecord "1" --> "0..*" CVALEarmarkRecord
CustomerSeedRecord --> PlacementPolicy : policy_version
PlacementPolicy "1" --> "1..*" RegionCatalogue : catalogue
SourceDestinationDRIndex "1" --> "0..*" ActivationRecord

```

```
QuotaPoolState "1" --> "0..*" ReservationState  
OperationRecord --> ReservationState : mutates
```

Note: the earlier UML summary's "two quota groups per region" is superseded by the **single governed QuotaPoolState** with logical earmarks (exception two-group topology only). [Decided]

Document Status: Draft for review · **Next:** ratify alongside the FDD; supersede the Production Readiness Review (archived). Production reliance on max-not-sum DR and single-pool consumer quota is gated on POC-011 / POC-001.