

ACRME Calculation Logic Reference — All Scenarios (v2.2)

Revision v2.2 — 27 August 2026.

This document supersedes the previous calculation logic reference. Key changes in v2.2:

- **DR sizing formula replaced:** fixed `dr_ratio_*` percentage approach (v1) is superseded by the **max-not-sum distributed-portion formula** (Requirements v2.1 Appendix A.6 / Appendix D).
- **Reservation target floor formula added:** `Allocated VM Count + Configured Buffer` (CAP-003), distinct from the forecast-based growth formula.
- **Four new scenarios added:** DR destination requirement (max-not-sum), standby activation, deployment readiness gate, and customer seed record.
- **EU cross-geo DR region corrected:** Belgium Central → **Switzerland North** throughout.
- **Quota grouping model updated (QUA-004):** **single governed pool** per region/quota family (Prod + NonProd/CVAL + DR together) is now the **primary model**; Prod/DR protected by logical earmarks. The two-group topology is retained only as a narrow governance/Azure-limit exception (Scenario 8/9 updated accordingly).
- **Scenario 15** (fixed DR ratio parameters) retained for reference but marked **superseded**.

This document consolidates **every calculation logic** used by the Azure Capacity Reservation Management Engine (ACRME), organized by the scenario in which it fires. Each formula carries an evidence tag:

- `[Documented]` — backed by Microsoft Learn or Azure platform documentation.
- `[Decided]` — an approved design decision recorded in the ADR Decision Log.
- `[Derived]` — a logical consequence of the design; requires POC validation where noted.
- `[Assumed]` — a policy default or working hypothesis, tunable and not yet empirically validated.

Sources: Requirements Baseline v2.1/v2.2 (Appendix A-D); ADR-001 through ADR-004 v1.2; `acrme_production_readiness_review_and_architecture.md` (PRR §26-§32).

Convention. `vCPU` = `vCPU_per_instance` for the SKU family. `CVAL` (Customer Validation) and `NonProd` are interchangeable; scoring identifiers keep the `NonProd` spelling for engineering consistency. All raw ratios are clamped with `Clamp(x) = max(0, min(1, x))` before weighting.

Index of Scenarios

#	Scenario	Core calculation	Status
1	Prod region derivation — customer picks a geography (exception path; explicit approval + customer acknowledgement required)	<code>argmax(PS_Prod)</code> over Standard regions	Current
2	Prod region validation — customer supplies the exact region (default input; AC-RME validates, does not derive)	HC-1..HC-10 gate + PS_Prod post-validation	Current
3	Restricted region request	Exception workflow (no scoring)	Current
4	Middle East three-region deployment	<code>argmax(PS_Prod)</code> in-geo + Switzerland North cross-geo DR	Updated
5	CVAL / NonProd region selection	<code>argmax(PS_NonProd)</code>	Current
6	DR region selection	<code>argmax(PS_DR)</code>	Current
7	Hard-constraint eligibility gate	HC-3, HC-6, HC-7 arithmetic	Current
8	Quota group sizing	Prod & NonProd+DR group budgets	Current
9	DR floor accounting	<code>DR_Floor_vCPU</code> , effective ceiling, headroom	Updated
10	Environment quota scoring	<code>Quota_Score_{Prod,NonProd,DR}</code>	Current
11	Capacity forecast / CR sizing	<code>Forecast_Quantity</code>	Current
12	Auto-increase trigger	Utilisation thresholds + debounce	Current
13			Current

#	Scenario	Core calculation	Status
	Emergency transfer & tier escalation	Tier 1/2/3, quota-neutral transfer	
14	Scaling & API-budget model	CRG cardinality, request-per-cycle	Current
15	DR ratio parameters (fixed %)	dr_ratio_min/max/target	Superseded by Scenario 17
16	Reservation target & reconciliation floor	Allocated VM Count + Buffer (CAP-003)	New
17	DR destination requirement — max-not-sum	MAX(source portions per destination) (A.6)	New
18	Standby activation (DR-019)	associated → allocated staged acquisition	New
19	Deployment readiness gate	READY / QUOTA_DEFICIT / STALE_STATE / ... (RDY-002)	New
20	Customer placement seed record	First-placement persistence + reuse policy	New

Shared Building Blocks (used by several scenarios)

Default scoring weights [Assumed]

```

alpha (α) = 0.30 # capacity / headroom signal
beta (β) = 0.20 # quota headroom signal
gamma (γ) = 0.25 # capacity-weighted distribution
delta (δ) = 0.15 # DR-coverage / overflow-integrity signal
epsilon(ε)= 0.10 # zone diversity
Clamp(x) = max(0, min(1, x)) # every component clamped to [0,1] before weighting

```

All five weights are tunable policy constants stored in `PlacementPolicy` (config-as-code, versioned). They sum to exactly 1.0 and apply across all three environment scoring formulas. [Assumed]

Named regional-snapshot quantities [Decided]

```
nonprod_crg.effective_free = nonprod_crg.free_slots - dr_overflow_reserve
dr_crg.coverage_ratio      = dr_crg.quantity / potential_dr_demand
potential_dr_demand(region) =  $\sum$  prod_allocated for all customers whose dr_region = re-
gion
```

Key formula disambiguation

Formula	Purpose	Source
Target Reserved Capacity = Allocated + Buffer	Continuous reconciliation floor (Scenario 16)	CAP-003
Forecast_Quantity = ceil(Peak × (1+Growth) + DR_Buffer)	Proactive growth ahead of demand (Scenario 11)	ADR-004
Destination DR Require- ment(d) = MAX(source por- tions)	DR standby sizing (Scenario 17)	A.6 / DR-017
dr_ratio_* constants	Superseded (Scenario 15) — configuration reference only	PRR §32

Scenario 1 — Prod Region Derivation (Customer Chooses a Geography)

Trigger: customer supplies a geography (e.g. “Europe”), not a specific region.

Status: Exception path — not the default (PLC-002). Requires explicit exception approval and customer acknowledgement that the **derived production region becomes fixed until an approved migration**

changes the seed. Cannot be invoked without a recorded exception reference. [Decided]

Logic:

1. Build candidate set = **Standard Capacity Regions** in the chosen geography. Restricted regions are excluded before scoring. [Decided]
2. Score every candidate with PS_Prod from the versioned regional snapshot.
3. Select derived Prod anchor = $\text{argmax}(\text{PS_Prod})$.
4. **Tie-break:** deterministic — first region in the Standard Capacity Region list order for that geography; also the cold-start default when no snapshot exists. [Derived]
5. Run sequential CVAL → DR selection (Scenarios 5 & 6) on the same Standard region pool.

PS_Prod(r) [Decided]

```

PS_Prod(r) =
  0.30 × Clamp(nonprod_crg.effective_free / prod_crg.quantity)    ← α: NonProd
headroom signal
  + 0.20 × Clamp(prod_crg.quota_headroom / prod_crg.quota_limit)  ← β: Prod quota
headroom
  + 0.25 × Clamp(1 - prod_customer_count / total_customers)      ← γ: distribution
fairness
  + 0.15 × Clamp(dr_crg.coverage_ratio)                          ← δ: DR readiness
signal
  + 0.10 × Clamp(az_count / 3)                                    ← ε: zone di-
versity

```

PS_Prod is **dual-purpose**: derives the Prod region in Scenario 1 and is reused for post-selection validation in Scenario 2. Every candidate score and the policy version are written to the **OperationRecord** for deterministic replay. **[Derived]**

Scenario 2 — Prod Region Validation (Customer Supplies a Specific Region)

Trigger: customer names an exact Azure region — the **default** input mode (PLC-001). ACRME **validates** the supplied region against HC-1..HC-10 and placement-policy rules; it does **not** derive a region from geography. **[Decided]**

Logic:

- **Standard Capacity Region:** validate against HC-1..HC-10 (Scenario 7). If eligible, it becomes the Prod anchor directly. **PS_Prod** computed for post-selection validation/audit only, not for selection.
- **Restricted Capacity Region:** do not score — route to the Exception Deployment Workflow (Scenario 3).

There is no **argmax** here. The customer's choice is authoritative once it passes the hard constraints. The first approved placement writes the **Customer Seed Record** (Scenario 20). **[Decided]**

Scenario 3 — Restricted Region Request (Exception Path)

Trigger: the requested or only feasible region is a **Restricted Capacity Region**.

Logic: no scoring formula is evaluated. The request diverts to the manual Exception Deployment Workflow (governance approval, EC-1..EC-4, VR-1..VR-11, HC-9/HC-10 controls). Restricted regions never enter the scoring pipeline. **[Decided]**

Exception conditions:

Code	Condition
EC-1	Customer requires the specific restricted region by contract
EC-2	No Standard region in the geography can supply sufficient capacity
EC-3	Data-residency requirement binds to the restricted region
EC-4	Scenario 2 input — restricted regions can never be engine-derived

Scenario 4 — Middle East Three-Region Deployment (Updated — Switzerland North)

Trigger: a Middle East deployment requiring three environments.

Logic:

1. Candidate in-geo Standard regions = { Saudi Arabia Central, UAE North } (both Standard)
2. Score both with PS_Prod.
3. Prod = $\text{argmax}(\text{PS_Prod})$ over the two.
4. CVAL = the remaining in-geo region (deterministic — only one candidate left).
5. DR = cross-geo extension region: Switzerland North (Europe)
because no third in-geo Standard region exists to satisfy region separation.

v2.2 correction (REG-002): The cross-geo DR extension region was incorrectly cited as “Belgium Central” in earlier material. The authoritative placement configuration specifies **Switzerland North** as the cross-geo extension for Middle East. All references to Belgium Central are superseded by Switzerland North.

Middle East is subject to the `DR_NOT_OFFERED` policy flag (DR-014) where legal/data-sovereignty prevents an acceptable DR design. When `DR_NOT_OFFERED = true` for the geography, Step 5 above is bypassed and the seed record records `dr_region = NOT_OFFERED`. [Decided]

Cross-geo extension constraints from ADR-001 apply to the DR region (Switzerland North). [Decided]

Scenario 5 — CVAL / NonProd Region Selection

Trigger: sequential step after the Prod anchor is fixed. Selects $\text{argmax}(\text{PS_NonProd})$ over eligible Standard regions.

PS_NonProd(r) — design-of-record [Decided]

```

PS_NonProd(r) =
  0.30 × Clamp(nonprod_crg.effective_free / nonprod_crg.quantity)    ← α: effective NonProd headroom
+ 0.20 × Clamp(nonprod_crg.quota_headroom / nonprod_crg.quota_limit) ← β: NonProd quota headroom
+ 0.25 × Clamp(1 - nonprod_customer_count / total_customers)         ← γ: distribution fairness
+ 0.15 × Clamp(nonprod_crg.effective_free / nonprod_crg.quantity)     ← δ: overflow capacity health
+ 0.10 × Clamp(az_count / 3)                                           ← ε: zone diversity

```

Corrected pilot variant [Undocumented – architectural judgement]

The design-of-record above duplicates the same signal under α and δ (combined weight 0.45). The PRR's corrected pilot formula removes the duplication:

```

PS_NonProd = 0.35 × Capacity + 0.25 × Quota + 0.25 × Distribution
              + 0.05 × DR_Overflow_Integrity + 0.10 × Zones

Distribution = 1 - Region_Assigned_Demand / Total_Assigned_Demand

```

Both variants are recorded. The design-of-record formula is the approved baseline; the pilot variant is the reviewer-recommended refinement. Neither is empirically validated yet.

CVAL/DR co-location rule (PLC-010): a customer's CVAL and DR may co-locate in the same destination

region. When co-located, CVAL capacity earmarked for DR activation must **not** be double-counted as both

live CVAL headroom and available DR headroom. The `CVALEarmarkRecord` tracks this. [Decided]

Scenario 6 — DR Region Selection

Trigger: final sequential step. Selects `argmax(PS_DR)` over eligible Standard regions (or the cross-geo region for Middle East).

PS_DR(r) [Decided]

```

PS_DR(r) =
  0.30 × Clamp(dr_crg.free_slots / dr_crg.quantity)    ← α: DR CRG headroom
+ 0.20 × Clamp(dr_crg.quota_headroom / dr_crg.quota_limit) ← β: DR quota headroom
+ 0.25 × Clamp(1 - dr_customer_count / total_customers) ← γ: distribution fairness
+ 0.15 × min(1.0, dr_crg.coverage_ratio / dr_ratio_target) ← δ: coverage ratio health
+ 0.10 × Clamp(az_count / 3)                             ← ε: zone diversity

```


Note on 6: `dr_ratio_target` in the PS_DR formula is now a **configurable bootstrap reference** rather than a fixed 30–40% constant. See Scenario 17 for the authoritative DR sizing formula. [Decided]

Scenario 7 — Hard-Constraint Eligibility Gate (Arithmetic Constraints)

Before any region is scored (Scenario 1) or accepted (Scenario 2), it must pass the hard constraints.

HC-3 QUOTA_FLOOR (per environment) [Decided]

```
Prod:    Prod_Group_headroom(R) ≥ requested_vm_count × vCPU
NonProd: nonprod_headroom(R)    ≥ requested_vm_count × vCPU    # uses effective_nonprod_ceiling
DR:      dr_headroom(R)         ≥ target_dr_qty      × vCPU
```

HC-6 DR_COVERAGE_FLOOR [Decided]

```
Region ineligible for DR if:
(dr_crg.free_slots + nonprod_crg.effective_free) < customer_requested_dr_slots
```

HC-7 DR_FLOOR_INTEGRITY [Decided]

```
REJECT NonProd placement in region R if:
nonprod_quota_used(R) + (requested_vm_count × vCPU) > effective_nonprod_ceiling(R)
```

HC-3 checks raw group headroom; HC-7 additionally enforces floor compliance. Both must pass. On breach:

emit `DRFloorViolationDetected` (Critical severity).

Minimum-headroom policy defaults [Decided]

```
min_prod_quota_headroom_vcpu    = 20
min_nonprod_quota_headroom_vcpu = 20
min_dr_headroom_vcpu             = 16
```

Scenario 8 — Quota Group Sizing (per Region)

Primary model (QUA-004): one governed quota pool per region and quota family, covering Prod,

NonProd/CVAL, and DR together. Prod and DR are protected inside the shared pool by logical earmarks, not by physical group separation. [Decided]

```

Pool_Limit(region) =
    Prod_CRG_quantity      × vCPU × (1 + prod_growth_buffer)
  + NonProd_CRG_quantity  × vCPU × (1 + nonprod_growth_buffer)
  + DR_Earmark_vCPU(region)
  + emergency_transfer_headroom_vcpu

prod_growth_buffer          = 0.20
nonprod_growth_buffer       = 0.20
DR_Earmark_vCPU(region)     = Destination_DR_Requirement(region) × vCPU    [max-
not-sum, Scenario 17]
emergency_transfer_headroom_vcpu = max_emergency_transfer_qty × vCPU
max_emergency_transfer_qty   = potential_dr_demand × emergency_transfer_pct
emergency_transfer_pct       = 0.30

```

Because Prod, NonProd/CVAL, and DR share one pool, an emergency DR draw needs **no cross-group transfer**:

the DR orchestrator consumes `Pool_Headroom` first, then reclaimable NonProd above committed demand.

This is the manipulation-flexibility gain that motivates the single-pool model (avoids fragmenting quota and cuts quota-increase requests to Microsoft). [Decided]

Prod protection inside the shared pool (logical earmarks, enforced at allocation/reclamation time):

```

Prod_Reserved_Floor = Prod_Used_vCPU + Prod_Growth_Buffer_vCPU
Allocatable_NonProd = Pool_Limit - Prod_Reserved_Floor - DR_Earmark_vCPU - Non-
Prod_Used_vCPU
Pool_Headroom       = Pool_Limit - Pool_Used

```

`Prod_Reserved_Floor` and `DR_Earmark_vCPU` are never allocatable to NonProd; NonProd allocation fail-saves when `Allocatable_NonProd` reaches zero. [Decided]

Exception topology (only when Azure Quota Group limits or a mandatory Prod-isolation governance

boundary make one pool impossible) — two groups per region, recorded in the Decision Log:

```

Prod_Group_Limit(region) =
    Prod_CRG_quantity × vCPU × (1 + prod_growth_buffer)

NonProd_DR_Group_Limit(region) =
    NonProd_CRG_quantity × vCPU × (1 + nonprod_growth_buffer)
  + DR_CRG_quantity      × vCPU
  + emergency_transfer_headroom_vcpu

```

In the exception topology the `emergency_transfer_headroom_vcpu` term is what makes Tier 2 DR expansion

quota-neutral within the shared NonProd+DR group. In the single-pool model this term is subsumed by the

shared `Pool_Headroom`. [Decided]

Quota-as-governor principle (QUA-005): quota allocation caps deployable capacity per product, environment, subscription, region, and VM family. Teams must justify quota increases — unallocated

pooled

quota is not automatically distributed. [Decided]

Scenario 9 — DR Floor Accounting (Engine-Enforced Sub-Limit) — Updated

Azure quota pools/groups have no native intra-pool sub-reservation; the engine enforces the DR floor by

arithmetic (logical earmark) in both the single-pool and exception two-group topologies. [Decided]

Accounting formulas [Decided]

```
DR_Floor_vCPU(region)      = Destination_DR_Requirement(region) × vCPU
                             [v2.2: uses max-not-sum formula from Scenario 17]
                             [v1: was potential_dr_demand × vCPU × dr_ratio_max - su-
perseded]

# Single governed pool (primary model, QUA-004):
DR_Earmark_vCPU(region)    = DR_Floor_vCPU(region)          # never allocatable to NonProd
Allocatable_NonProd        = Pool_Limit - Prod_Reserved_Floor - DR_Earmark_vCPU - Non-
Prod_Used_vCPU
Pool_Headroom              = Pool_Limit - Pool_Used

# Exception two-group topology only:
Effective_NonProd_Ceiling = NonProd_DR_Group_Limit - DR_Floor_vCPU
NonProd_Headroom          = Effective_NonProd_Ceiling - NonProd_Used_vCPU
Group_Headroom            = Group_Limit - Group_Used

dr_crg.coverage_ratio      = dr_crg.quantity / Destination_DR_Requirement(region)
```

v2.2 change: DR_Floor_vCPU is no longer derived from potential_dr_demand × dr_ratio_max. It is now derived from Destination_DR_Requirement(region) — the max-not-sum value from Scenario 17.

This aligns the floor with the lean bootstrap model (DR-007) and the distributed DR topology (DR-016).

The fixed dr_ratio_max = 0.40 constant is retained in Scenario 15 for legacy reference only.

Dual-validation control [Decided]

A separate detector recomputes the DR floor from authoritative assignment data; any disagreement between

command-time and detector values disables automatic NonProd expansion until reconciled.

Worked topology example [Decided]

```
Prod group      = 128 vCPU  (32 × D4s_v3, 4 vCPU each)
NonProdDR group = 80 vCPU
DR floor        = 32 vCPU   (= Destination_DR_Requirement for this region)
Effective NonProd ceiling = 48 vCPU (80 - 32)
```

Scenario 10 — Environment Quota Scoring (β Component Detail)

The generic `Quota_Score` is dispatched by environment type. [Decided]

```
Quota_Score_Prod(r)    = prod_group_headroom_vcpu / prod_group_limit_vcpu
Quota_Score_NonProd(r) = nonprod_headroom_vcpu   / effective_nonprod_ceiling_vcpu
Quota_Score_DR(r)     = dr_headroom_vcpu         / dr_floor_vcpu
```

Semantics differ by type: a DR score of `1.0` means maximum expansion room (DR CRG at zero), not maximum quota consumed. Zero-denominator guards apply. [Decided]

Scenario 11 — Capacity Forecast / CR Sizing

Advisory forecast used to size CR quantity over a horizon — the **proactive growth path**, distinct from

the continuous reconciliation floor (Scenario 16). [Decided]

```
Forecast_Quantity = ceil( Forecast_Peak × (1 + Growth_Buffer) + DR_Buffer )

Forecast_Peak      = predicted peak allocated-VM demand in the horizon      [De-
rived]
Growth_Buffer      = policy % for forecast uncertainty                      [As-
sumed]
DR_Buffer          = extra units required by approved recovery policy        [As-
sumed]
Forecast_Horizon ∈ { 30, 60, 90 } days                                     [Assumed]
```

Lead-time alerting: when forecast demand approaches 80% of quota limit, emit

`ForecastApproachingQuotaLimit` with 14-day lead time. [Documented]

Forecast recommendations remain advisory until model accuracy and false-positive rates are measured.

The horizon and buffer values are policy percentages stored in `PlacementPolicy`, not fixed constants.

Scenario 12 — Auto-Increase Trigger (Steady-State Capacity Lifecycle)

When utilisation crosses a threshold, a `CapacityIncreaseRequest` is raised. [Decided]

```
Trigger auto-increase when utilisation ≥ threshold:
  dr_autoincrease_threshold      = 0.35
  prod_autoincrease_threshold    = 0.20
  nonprod_autoincrease_threshold = 0.20
```

Debounce cooldown = 30 min per (region + CRG type) after any trigger

Phase 1 requires operator approval. Auto-decrease excluded from Phase 1. The steady-state lifecycle runs **only** in `STEADY_STATE` engine mode (see ADR-003). `[Decided]`

Normative 10-step steady-state capacity lifecycle

1. Detect threshold crossing.
2. Re-read current CR, quota, sharing, and assignment state.
3. Create `CapacityIncreaseRequest`.
4. Calculate target quantity (`Forecast_Quantity` or `Allocated + Buffer` floor, whichever is higher).
5. Require operator approval (Phase 1).
6. Submit the quota action only if validated as required.
7. Wait for confirmed quota state — no assumed propagation SLA.
8. Update CR quantity.
9. Confirm the actual quantity.
10. Refresh the snapshot and close the request.

Scenario 13 — Emergency Transfer & Tier Escalation (DR Event Active)

During an active DR event (`DR_EVENT_ACTIVE` engine mode), emergency capacity is obtained by tier. `[Decided]`

Tier 1	<code>DirectExpansion</code>	– additive DR expansion using existing headroom. Automated.
Tier 2	<code>QuotaNeutralTransfer</code>	– reduce NonProd CR, expand DR within the SAME quota group. Policy-gated. Quota-neutral (see below).
Tier 3	<code>DestructiveTransfer</code>	– changes VM associations. Dual-approval + elevated RBAC. BLOCKED in Phase 1.

Quota-neutral math for Tier 2 `[Derived – P0C-31/32]`

Because NonProd CR and DR CR draw from the SAME NonProd+DR quota group:
 NonProd CR reduction → releases $q \times \text{vCPU}$ to the GROUP pool
 DR CR expansion → consumes $q \times \text{vCPU}$ from the SAME GROUP pool
 Net group headroom change ≈ 0 → no Azure quota-increase request needed
 Tier RTO gated only by ARM operation time (minutes), not quota approval (hours)

$$\text{max_emergency_transfer_qty} = \text{potential_dr_demand} \times \text{emergency_transfer_pct} \quad (\text{pct} = 0.30)$$

Escalation order: `Tier 1 headroom available?` → `Tier 2 approved?` → `Tier 3 allowed?` → manual.

Scenario 14 — Scaling & API-Budget Model

Used to size the estate and stay within Azure Resource Manager throttling budgets. `[Derived]`

```

CRG_lower_bound = 3 × R                                # 3 CRGs (Prod/NonProd/DR) per region
CRG_total       = R × Eclass × Z × SKUset × IsolationFactor

Requests_per_cycle = CRG_reads + CR_reads + sharing_reads + quota_reads + associ-
ation_reads
Average_requests_per_second = Requests_per_cycle / 300      # 5-minute (300 s)
cycle

```

Documented ARM throttle baselines [Documented]

```

Compute RP read limit = 250 requests / 5 min / subscription
Compute RP write limit = 1,200 requests / hour / subscription

```

The adaptive throttle manager initialises with these baselines and adapts to observed `429/Retry-After`.

Reconciliation cadence [Assumed]

```

Reconciliation loop target      = 6 min (configurable; production interval TBD)
Critical targeted reconcile P95 < 2 min
Stable-resource state age P95   < 15 min
Drift detection                 ≤ 2 reconciliation cycles

```

Scenario 15 — DR Ratio Parameters (Fixed %) — SUPERSEDED

Status: Superseded by Scenario 17 (max-not-sum). These constants are retained for configuration reference and legacy comparison only. The fixed-percentage DR sizing model (30–40% of Prod) was rejected in Requirements v2.0/v2.1 because it produces idle reserves estimated at **\$1.5M–\$5M/year** at platform scale — directly contradicting the cost-reduction mandate. See Appendix D of the Requirements Baseline for the full derivation.

```

dr_ratio_min    = 0.30      [superseded – was HC-6 lower bound]
dr_ratio_max    = 0.40      [superseded – was DR_Floor_vCPU multiplier]
dr_ratio_target = [0.30, 0.40] [superseded – now: configurable bootstrap per
product/workload]
emergency_transfer_pct = 0.30 [still current – Scenario 13]
prod_growth_buffer   = 0.20  [still current – Scenario 8]
nonprod_growth_buffer = 0.20  [still current – Scenario 8]

```

The `SUM` override (C-11 in the Configurable Items Register) remains available for specific customers or geographies that contractually require protection against concurrent regional failures. All other scopes use the max-not-sum formula (Scenario 17).

Scenario 16 — Reservation Target & Reconciliation Floor (CAP-003) — New

Trigger: every reconciliation cycle compares current reserved quantity against this floor.

This is the **continuous floor formula** — distinct from the forecast-based growth formula (Scenario 11).

A.1 Reservation target (CAP-003) [Decided]

Target Reserved Capacity = Allocated VM Count + Configured Buffer

- **Allocated VM Count** = VMs in **running** / **allocated** state consuming compute capacity.
- **Configured Buffer** = policy-defined headroom above allocated demand. Tunable per product/region/env/SKU. Not hard-coded.
- **Associated-but-deallocated VMs** do not force reservation retention (CAP-004). They are reported separately so teams understand restart risk, but do not add to the target.

A.2 Reservation headroom

Reservation Headroom = Reserved Quantity - Allocated VM Count

A.3 Reservation deficit

Reservation Deficit = $\max(0, \text{Target Reserved Capacity} - \text{Reserved Quantity})$

Reconciliation decision logic [Decided]

```
IF Reserved Quantity < Target Reserved Capacity:
    RAISE reservation (scale-up) toward target, subject to Azure availability.
    IF Azure cannot supply: hold current state, RAISE alert, expose buffer deficit.

IF Reserved Quantity > Target Reserved Capacity (consistently, after minimum-hold interval):
    LOWER reservation toward target – ONLY when all guards pass:
        - current_reserved > allocated_count           (never below allocated)
        - not within DR protection window
        - no active maintenance exclusion
        - cost policy permits reduction
```

Zero-capacity and no-auto-delete policy (CAP-009/010) [Decided]

Reduce an unused managed reservation to ZERO; do NOT delete the reservation object. Normal reconciliation NEVER deletes CRGs or reservation definitions. Deletion uses a separate approved decommissioning workflow only.

Scenario 17 — DR Destination Requirement — Max-Not-Sum (A.6/DR-017) — New

This scenario replaces the fixed `dr_ratio_*` sizing in Scenario 15.

The single-failure assumption (DR-001) [Decided]

The programme plans for failure of **one** production region within a geography at a time. Simultaneous multi-region failure is outside the default guaranteed model.

A.6 DR destination requirement — corrected formula [Decided]

```
Destination_DR_Requirement(d)
  = MAX over each non-concurrent source region s protected by d (
    Workload_Portion(s → d)
  )
```

Where:

- `d` = destination region hosting DR standby capacity
- `s` = source region whose production customers have DR in `d`
- `Workload_Portion(s → d)` = the quantity of VMs/vCPUs from source `s` that are assigned to land in destination `d` on failover (from the source→destination DR index, DR-018)

Rationale: because only one region fails at a time, destination `d` never simultaneously hosts failover from more than one source. It therefore needs standby capacity for its **largest** protected source only — never the sum of all sources. The standby slots are **shared / overcommitted** across mutually exclusive failure events. [Decided]

A.7 DR capacity gap [Decided]

```
DR_Capacity_Gap(d) = max(0,
  Destination_DR_Requirement(d) - Usable_Destination_Capacity(d)
)
```

`Usable_Destination_Capacity(d)` may include approved bootstrap headroom, available CRG reservations, releasable CVAL capacity (after earmark check, Scenario 9/PLC-010), and capacity acquired through approved sharing or expansion.

A.8 Shared-DR overcommit ratio (informational) [Derived]

```
Overcommit_Ratio(d) = SUM(Workload_Portion(s → d) for all s) / MAX(Workload_Portion(s
→ d))
```

A ratio > 1 quantifies: (a) the capacity and cost saved by max-not-sum sizing, and (b) the exposure if the single-failure assumption is ever violated. Required input for POC-011 risk sign-off.

Observability requirement (OBS-002) [Decided]

Alert when a destination's usable capacity falls below `Destination_DR_Requirement(d)` — the max protected source — not below the sum. Dashboard must show per-destination max-source coverage (OBS-004).

Conservative SUM override (C-11) [Decided]

`Destination_DR_Requirement_Conservative(d) = SUM(Workload_Portion(s → d) for all s)`

Used only where a customer or geography contract explicitly requires protection against concurrent regional failures. Configured per-scope in `PlacementPolicy`; no code change required.

Worked example (from Requirements Appendix D) [Decided]

Region 2 holds DR standby for customers from Region 1 and Region 3:

Source protected by R2	Failover portion in R2
Region 1 (Cust1 + Cust5)	120 cores
Region 3 (Cust6)	80 cores

SUM sizing (old): $\text{Requirement}(R2) = 120 + 80 = 200$ cores ← over-provisions; super-seded

MAX sizing (new): $\text{Requirement}(R2) = \max(120, 80) = 120$ cores

Saving at R2: $200 \rightarrow 120 = 80$ cores (40%) removed with no loss of protection.

$\text{Overcommit_Ratio}(R2) = 200 / 120 \approx 1.67$

→ R2's shared standby is 1.67× oversubscribed.

→ Platform-wide extrapolation: ~40–50% standby reduction ≈ \$0.6M–\$2.5M/year saved.

Source→Destination DR Index entity (DR-018) [Decided]

The max-not-sum formula depends on an authoritative bidirectional mapping — a required state-store entity:

```
SourceDestinationDRIndex {
  source_region      : string
  destination_region : string
  customer_realm_id  : string
  standby_instance_set : list<InstanceRef>
  sku_family         : string
  quantity_vcpu      : int
  activation_state    : STANDBY | ACTIVATING | ACTIVE | FAILBACK_PENDING
  last_updated       : datetime
  policy_version      : string
}
```

This index is the reverse view of the customer seed record (Scenario 20). On a regional failure, the engine queries this index to determine exactly which standby sets to activate and where. [Decided]

Scenario 18 — Standby Activation (DR-019) — New

Trigger: authorised DR declaration; engine mode transitions to `DR_EVENT_ACTIVE`.

Per-customer activation workflow [Decided]

Each customer's DR instances transition `associated` → `allocated` (`standby` → `active`) in approved business-priority waves:

```
Priority Wave 0 (P0): Platform control planes and recovery orchestration infrastruc-
ture.
Priority Wave 1 (P-1): Highest-priority customer workloads.
Priority Wave N: Remaining customers in business-priority order.
```

DR-006 staged capacity acquisition per wave:

```
Stage 1: Use approved bootstrap / pre-staged headroom (zero Azure wait time).
Stage 2: Allocate available reservation + quota already in the destination.
Stage 3: Shut down / disassociate eligible CVAL workloads (after DR-010 authorised
trigger).
Stage 4: Share or reassign reservations within supported region/zone boundaries.
Stage 5: Allocate pooled quota to DR subscriptions.
Stage 6: Request additional Azure quota/capacity where required.
Stage 7: Report unrecoverable capacity gaps.
```

Stages 1-2 start immediately from bootstrap headroom — P0/P-1 customers begin recovery without waiting

on Azure deallocation/disassociation APIs (which are throttled during regional events). [Decided]

Activation state tracking [Decided]

```
ActivationRecord {
    customer_realm_id : string
    source_region      : string
    destination_region : string
    priority_wave      : int
    activation_state    : PENDING | ACTIVATING | ACTIVE | FAILED | FAILBACK_PENDING
    stage_reached       : int          # 1-7 per DR-006
    capacity_gap_vcpu   : int          # 0 if fully provisioned
    activated_at        : datetime
    approved_by         : string
    incident_id         : string
}
```

Activation state is tracked per customer, auditable, and reversible on failback (DR-013). [Decided]

Failback reversal [Decided]

On failback: `ACTIVE` → `FAILBACK_PENDING` → `STANDBY` in reverse priority order. The seed record and DR index are preserved; no re-derivation of placement. History, audit trail, and DR capacity state are retained through the full failback cycle.

Scenario 19 — Deployment Readiness Gate (RDY-002) — New

Trigger: every AEP-triggered deployment before the first environment is provisioned.

Machine-readable readiness states [Decided]

READY	All gates pass; deployment may proceed.
READY_WITH_RISK	All gates pass but one or more advisory risks detected (e.g. buffer below target).
QUOTA_DEFICIT	Required deployment quota > available quota in the deploying subscription.
RESERVATION_DEFICIT	Reservation exists but quantity insufficient; approved over-allocation not in effect.
CAPACITY_UNAVAILABLE	Azure cannot supply the requested SKU/zone at this time.
STALE_STATE	Capacity/quota snapshot exceeds max-age threshold; cannot drive a safe placement.
POLICY_BLOCKED	A governance policy (hard constraint, exception, approval) blocks the deployment.
VALIDATION_REQUIRED	One or more validation steps (e.g., POC-gated behaviours) not yet confirmed.

Readiness gate logic (RDY-001) [Decided]

All of the following must pass for `READY` :

1. `target_region` $\in$ `approved_region_catalogue` (REG-001)
2. `az_count` ≥ 1 in `target_region` (CAP-011)
3. `sku` $\in$ `supported_skus(target_region)`
4. `reservation_policy` known for (`subscription`, `region`, `zone`, `sku`)
5. `reservation` exists (if required by policy)
6. `reserved_quantity` $\geq$ `requested_count` OR `approved_over_allocation` active
7. `consumer_subscription_quota` $\geq$ `requested_units` (QUA-013 – POC-gated)
8. `snapshot_age` $\leq$ `max_snapshot_age_seconds` (RDY-004)
9. no blocking policy exception or hard constraint violation

Staleness check (RDY-004) [Decided]

```
IF now() - snapshot.collected_at > max_snapshot_age:
    RETURN STALE_STATE
    ACTION: refresh synchronously OR block with STALE_STATE – never drive placement
    from stale data
```

Quota deficit formula (A.5) [Decided]

```
Quota_Deficit = max(0, Requested_Deployment_Units - Available_Quota)
Available_Quota = Assigned_Regional_VM_Family_Quota - Current_Regional_VM_Family_Usage (A.4)
```

Scenario 20 — Customer Placement Seed Record (PLC-003/004/005) — New

Trigger: first approved production-region placement for a customer in a geography.

Seed record creation [Decided]

```
CustomerSeedRecord {
  customer_realm_id      : string      # authoritative customer/realm identifier
  geography              : string      # e.g. "NorthAmerica", "Europe"
  production_region      : string      # exact Azure region name
  cval_region            : string      # exact Azure region name
  dr_region              : string | "NOT_OFFERED"
  products_covered       : list<string> # all products seeded by this record
  decision_timestamp     : datetime
  policy_version         : string      # PlacementPolicy version used
  capacity_snapshot_ref  : string      # snapshot ID that drove the decision
  exception_ref          : string | null # approval ID if Restricted region or geo-
exception
  approval_metadata      : ApprovalRecord
  input_mode             : "SPECIFIC_REGION" | "GEOGRAPHY_EXCEPTION"
}
```

Seed reuse policy (PLC-004) [Decided]

```
IF CustomerSeedRecord exists for (customer_realm_id, geography):
  READ seed record – do NOT re-invoke the placement engine.
  Subsequent products and environments use the seeded production/CVAL/DR regions.
  Engine is NOT re-run per product once seeded.
```

Controlled seed change (PLC-005) [Decided]

```
IF seed change requested:
  REQUIRE approved migration/exception workflow.
  REQUIRE impact analysis across all products covered by the seed.
  PROHIBIT automatic regeneration on: upgrades, rebuilds, routine deployments.
```

A. Core Formula Reference (Appendix A — Requirements v2.1)

Label	Formula	Source
A.1	Target Reserved Capacity = Allocated VM Count + Con- figured Buffer	CAP-003
A.2	Reservation Headroom = Re- served Quantity - Allocated VM Count	CAP-003
A.3	Reservation Deficit = max(0, Target Reserved Ca- pacity - Reserved Quantity)	CAP-003
A.4	Available Quota = Assigned Regional VM-Family Quota - Current Usage	QUA-002
A.5	Quota Deficit = max(0, Re- quested Deployment Units - Available Quota)	QUA-007
A.6	Destina- tion_DR_Requirement(d) = MAX over s (Work- load_Portion(s → d))	DR-017
A.7	DR_Capacity_Gap(d) = max(0, DR_Requirement(d) - Us- able_Capacity(d))	DR-017
A.8	Overcommit_Ratio(d) = SUM(portions) / MAX(portions)	DR-017

B. Consolidated Policy-Constant Table (v2.2)

Constant	Value	Used in Scenario(s)	Status
alpha / beta / gamma / delta / epsilon	0.30 / 0.20 / 0.25 / 0.15 / 0.10	1, 5, 6	Current
prod_growth_buffer	0.20	8	Current
non-prod_growth_buffer	0.20	8	Current
emergency_transfer_pct	0.30	8, 13	Current
min_prod_quota_headroom_vcpu	20	7	Current
min_nonprod_quota_headroom_vcpu	20	7	Current
min_dr_headroom_vcpu	16	7	Current
dr_autoincrease_threshold	0.35	12	Current
prod/non-prod_autoincrease_threshold	0.20	12	Current
Debounce cooldown	30 min	12	Current
Reconciliation loop target	6 min (configurable)	14	Current
ARM read / write baseline	250 per 5 min / 1,200 per hour	14	Current
Forecast_Horizon	30 / 60 / 90 days	11	Current
dr_ratio_min	0.30	15	Superseded
dr_ratio_max	0.40	15	Superseded
dr_ratio_target	[0.30, 0.40]	15	Superseded
DR bootstrap target	Configurable per product/workload — no fixed %	17	Replaces ratio

Constant	Value	Used in Scenario(s)	Status
Max-not-sum default	MAX(source portions)	17	Current
SUM override (C-11)	SUM(source portions) — per-scope opt-in	17	Current
EU cross-geo DR extension region	Switzerland North	4	Updated (was Belgium Central)

C. Evidence & Maturity Summary

Formula	Evidence level	Validation required
PS_Prod , PS_NonProd , PS_DR (scoring)	[Decided] — design-of-record	Shadow/recommendation mode until empirically validated
HC-3, HC-6, HC-7 (hard constraints)	[Decided]	Ready for Phase 1
DR floor accounting (Scenario 9)	[Decided]	Updated to use max-not-sum input
Tier 2 quota-neutral math (Scenario 13)	[Derived]	POC-31/POC-32 required
Max-not-sum DR sizing (Scenario 17)	[Decided]	POC-011 required before production dependency
Consumer-subscription quota (QUA-013)	[Assumed]	POC-001 — top technical unknown
Reservation target floor (Scenario 16)	[Decided]	Ready for Phase 1
Standby activation staging (Scenario 18)	[Decided]	Dependent on POC-005 (VM state semantics)
Deployment readiness gate (Scenario 19)	[Decided]	Ready for Phase 1
Customer seed record (Scenario 20)	[Decided]	Ready for Phase 3

All constants are policy defaults stored in `PlacementPolicy` (config-as-code, versioned). Tuning any constant requires updating the config; no code change is needed. The `dr_ratio_*` constants are retained in the codebase as fallback references for the SUM override (C-11) only.

Document version 2.2 — 27 August 2026. Next review: upon POC-001 / POC-011 results.