

ACRME — Technical Design Document (TDD)

Title	Azure Capacity Reservation Management Engine (ACRME) — Technical Design Document
Version	1.0 (net-new)
Date	2 September 2026
Status	Draft for review — supersedes the Production Readiness Review as the technical design of record
Baseline	Azure Capacity & Quota Management — Consolidated Requirements Baseline v2.2 (27 Aug 2026)
Owner	Vishnuvardhan Reddy · Principal Cloud Architect
Audience	Engineering, SRE/platform operations, security, POC leads
Companion	Functional Design Document (<code>acrme_functional_design_document.md</code>)
Decision records	ADR-001 (region selection), ADR-002 (quota & capacity, single-pool), ADR-003 (capacity during DR), ADR-004 (forecast & increase), ADR-005 (distributed DR reference model) — all v2.2

Purpose. This TDD describes **how** ACRME is built — components, runtime, topology, data, state, algorithms, interfaces, security, observability, NFRs, integration, and POC gating — traceable to Requirements Baseline v2.2 and to the companion FDD (which owns the *what*). This document is **self-contained**: all normative algorithms, formulas, schemas, enumerations, and constants are in-lined, not referenced externally.

Reconciliation note (v2.2). This TDD implements the confirmed v2.2 decisions: **single governed quota pool** as the primary technical model with logical earmarks for Prod/DR protection (QUA-004, ADR-002); **Switzerland North** as the EU cross-geo DR extension region (REG-002) — pre-configured but **conditional and currently inactive** for the Middle East, whose current legal position is `DR_NOT_OFFERED` pending DEC-001 (see Section 2.2 constraints); **max-not-sum** destination DR sizing with an authoritative `SourceDestinationDRIndex` (DR-016/017/018, ADR-005); **exact-production-region-first** validation with a

governed `CustomerSeedRecord` (PLC-001..005, ADR-001); **five-state engine machine** and **standby activation waves** (DR-019, ADR-003).

1. Introduction

1.1 Scope

This document specifies the technical design of ACRME: the control engine that guarantees every managed Azure deployment has **both** reserved physical capacity and deployable VM-family quota — in the correct region/zone/SKU — before it proceeds, while minimising idle cost through lean distributed DR and continuous reconciliation. It covers component architecture, runtime and deployment, data and state, the placement/scoring/DR algorithms, the API surface, security, observability, non-functional behaviour, integration, and the POC dependencies that gate production reliance.

1.2 Relationship to the FDD

The FDD (`acrme_functional_design_document.md`) defines **what** ACRME does — capabilities, flows, states, and rules — traceable to Baseline v2.2. This TDD defines **how** those capabilities are realised. Where the FDD names a behaviour (e.g. “single governed pool”, “max-not-sum DR”, “readiness states”), this TDD gives the component, schema, algorithm, and interface that implements it. Section references to the FDD use the form (*FDD Section 4.x*).

1.3 Decision records referenced

ADR	Title	Key technical mandate used here
ADR-001	Region selection & customer placement	Exact-region-first validation; seed record; staged placement pipeline; input modes
ADR-002	Quota & capacity management	Single governed pool (primary) ; logical earmarks; exception two-group topology
ADR-003	Capacity management during DR	Five-state engine machine; three-tier emergency transfer; quota-neutral Tier 2

ADR	Title	Key technical mandate used here
ADR-004	Forecast & increase of capacity/quota	Reservation floor vs forecast growth; auto-increase triggers; 10-step lifecycle
ADR-005	Distributed DR reference model	<code>SourceDestinationDRIndex</code> ; max-not-sum; overcommit ratio/gap; Section 12A worked model

1.4 Evidence tags

Every non-trivial design assertion carries an evidence tag consistent with the requirements corpus: `[Documented]` (Microsoft Learn / Azure platform), `[Decided]` (approved ADR decision), `[Derived]` (logical consequence; POC where noted), `[Assumed]` (tunable policy default, not yet empirically validated).

2. Architecture Overview

2.1 Design principles (baseline Section 20)

1. **Readiness is capacity quota freshness.** No deployment is “ready” on a reservation alone; deployable consumer quota and non-stale state are equally required (RDY-003, NFR-002). `[Decided]`
2. **Continuous reconciliation of desired vs actual.** A reservation floor (`Allocated + Buffer`, CAP-003) is enforced every cycle; forecast growth is a separate proactive path (ADR-004). `[Decided]`
3. **Lean, distributed, reciprocal DR.** No dedicated DR regions; each region hosts Prod + CVAL + DR standby for other regions; standby is sized max-not-sum and shared across mutually exclusive failures (DR-016/017, ADR-005). `[Decided]`
4. **Single governed quota pool.** Collect all eligible per-region quota into one pool per region/quota family and allocate on demand; protect Prod/DR by logical earmark, minimising quota-increase requests to Microsoft (QUA-004, ADR-002). `[Decided]`
5. **Config-as-code, deterministic, replayable.** All policy (weights, buffers, thresholds, catalogue) is versioned; every decision records inputs, policy version, and snapshot ref for deterministic replay. `[Decided]`
6. **Fail-safe and idempotent.** Stale state blocks placement; all mutations are idempotent with operation handles and optimistic concurrency (NFR-002/007). `[Decided]`

2.2 Constraints (baseline Section 21)

- **Middle East DR is DR_NOT_OFFERED (current legal position, DEC-001 pending).** The Middle East programme is legal-owned and serves government/medical customers under data-residency/sovereignty laws; cross-border DR cannot meet residency requirements, so **no DR is offered there today** (baseline Section 2, Section 6, DR-014, DEC-001). Middle East **production** may still be placed and governed; the placement/DR engine emits `dr_region = NOT_OFFERED` for those geographies and creates **no** `SourceDestinationDRIndex` entry, **no** DR earmark, and **no** cross-geo activation. The Switzerland North cross-geo path (REG-002/HC-10) is **pre-configured but inactive** for the Middle East; it activates **only** if Legal clears DEC-001, via a config flip (`dr_not_offered["Middle East"] = false`) with **no code change**. `DR_NOT_OFFERED` is evaluated **before** any cross-geo extension. [Documented]
- **Azure Quota Group / groupQuotas maturity** gates single-pool enforcement (DEP-001); the two-group topology is the sanctioned fallback. [Documented]
- **No native intra-pool sub-reservation** in Azure quota groups — the DR earmark and Prod floor are engine-enforced arithmetic, not platform primitives (Scenario 9). [Documented]
- **ARM throttling budgets** bound reconciliation cadence and estate size (Compute RP: 250 reads / 5 min, 1,200 writes / hour per subscription). [Documented]
- **Phase 1 guardrails:** destructive VM-association transfers (Tier 3) are blocked; auto-decrease is operator-gated; forecast is advisory. [Decided]
- **Consumer-quota-under-shared-reservation (POC-001)** is the top unknown; readiness must validate consumer quota until proven. [Assumed]

3. Logical Component Architecture (Section 23)

ACRME is a set of cooperating services around a versioned state store and a policy (config-as-code) service. Components are deployable as modules of a single control application; boundaries are drawn for testability and independent scaling.

Component	Responsibility	Key requirements
Inventory Collector	Polls Azure (CRGs, reservations, quota/groupQuotas, usage, associations, zones) into freshness-stamped snapshots.	DAT-001..004, CAP-001/002, NFR-002
State Reconciler	Every cycle compares desired vs actual; enforces Allocated + Buffer floor; guarded scale-down; never deletes.	CAP-003/004/009/010, OPS-002
Placement & Scoring Engine	Validates exact region (or derives on exception); computes PS_Prod/PS_NonProd/PS_DR; writes seed record.	PLC-001..010, REG-001..003
Quota-Pool Manager	Maintains the single governed pool; computes earmarks; allocates on demand; enforces Prod floor & DR earmark; exception two-group mode.	QUA-001..014, ADR-002
DR Orchestrator	Maintains SourceDestinationDRIndex ; sizes max-not-sum; runs standby activation waves & staged acquisition; failback.	DR-001..019, ADR-005
Readiness API	Combined capacity+quota+freshness verdict as a machine-readable state to AEP.	RDY-001..004, INT-001..007
State Store	Versioned, optimistic-concurrency document store for entities, snapshots, operation/audit records.	DAT-001..006, NFR-007

Component	Responsibility	Key requirements
Config / Scope-File Service	Versioned PlacementPolicy, region catalogue, scope files; decision-traceable.	GOV-003, DAT-005, REG-001
Forecast & Increase Service	Advisory Forecast_Quantity; auto-increase triggers; operator-gated capacity/quota increase.	CAP-005, ADR-004
Emergency Transfer Service	Tier 1/2/3 escalation during DR_EVENT_ACTIVE; quota-neutral Tier 2; Tier 3 blocked in Phase 1.	ADR-003, Section 13 baseline
Observability Emitter	Metrics, alerts, dashboards incl. per-destination max-source coverage & source destination map.	OBS-001..005
Governance & Audit	RBAC, managed identity, approval gates, break-glass, immutable audit with before/after + correlation IDs.	GOV-001..009

T1 — Logical component architecture

flowchart TB

```

    subgraph External
        AEP[AEP / Provisioning]
        OPS[Operators / DR coordinator]
        AZ[Azure control planes:<br/>Compute / Reservations / groupQuotas]
        OBS[Observability platform]
    end
    subgraph ACRME
        API[Readiness API]
        PE[Placement & Scoring Engine]
        REC[State Reconciler]
        QPM[Quota-Pool Manager]
        DRO[DR Orchestrator]
        FIS[Forecast & Increase Service]
        ETS[Emergency Transfer Service]
        INV[Inventory Collector]
    end

```

```

    CFG[Config / Scope-File Service]
    GOV[Governance & Audit]
    OE[Observability Emitter]
    STORE[(State Store<br/>versioned + optimistic concurrency)]
end
AEP -->|readiness query / mutations| API
OPS -->|approvals, DR declaration| API
API --> PE & QPM & DRO & REC
PE --> STORE
QPM --> STORE
DRO --> STORE
REC --> STORE
FIS --> STORE
ETS --> STORE
INV -->|snapshots| STORE
INV <--> AZ
QPM <--> AZ
REC <--> AZ
DRO <--> AZ
CFG --> PE & QPM & DRO & REC & FIS
GOV -.enforces.-> API
OE --> OBS
STORE --> OE

```

4. Runtime & Deployment (Section 23, CAP-006, PLC-009)

4.1 Execution model

- **Reconciliation job (container app):** a long-running scheduler executes the reconciliation loop on a **6-minute target interval (configurable, CAP-006)**; targeted/critical reconciles run out-of-band. It drives the Inventory Collector → State Reconciler → Quota-Pool Manager cycle and emits snapshots. [Decided]
- **Placement flow (function/event path, PLC-009):** onboarding and AEP readiness requests are handled synchronously against the latest snapshot; on stale state the request triggers a synchronous refresh or returns STALE_STATE. [Decided]
- **DR path:** DR declaration flips `engine_mode` to `DR_EVENT_ACTIVE`; the DR Orchestrator drives priority-wave activation out of the steady-state cadence. [Decided]

4.2 Cadence, SLOs and drift (Section 38, NFR)

Reconciliation loop target	= 6 min (configurable; production interval TBD)	[Assumed]
Critical targeted reconcile P95	< 2 min	[Assumed]
Stable-resource state age P95	< 15 min	[Assumed]
Drift detection	2 reconciliation cycles	[Assumed]
Snapshot max age (staleness gate)	= max_snapshot_age_seconds (config, RDY-004)	[Decided]

4.3 API surface (summary; detail in Section 11)

A readiness/placement API (query + idempotent mutations), an operations/approval API (increase approval, DR declaration, override, sharing activation), and a read-only audit/observability API. All mutating operations are idempotent and versioned (INT-004..006). [Decided]

T2 — Management group / subscription topology

flowchart TB

```
MG[Management Group root] --> MGPlat[Platform MG]
MG --> MGWork[Workload MGs]
MGPlat --> SubCtl[ACRME control subscription<br/>engine, state store, config]
MGPlat --> SubObs[Observability subscription]
MGWork --> SubProd[Prod subscriptions<br/>Prod CRGs + pool]
MGWork --> SubNonProd[NonProd/CVAL subscriptions<br/>NonProd CRGs]
MGWork --> SubDR[DR-hosting subscriptions<br/>DR standby CRGs]
SubCtl -. managed identity + RBAC .-> SubProd
SubCtl -. managed identity + RBAC .-> SubNonProd
SubCtl -. managed identity + RBAC .-> SubDR
```

5. Topology — Subscriptions, CRGs and Quota Groups (Section 24, Section 25, Section 26)

5.1 Management group / subscription model (Section 24)

ACRME runs from a platform control subscription with cross-subscription managed-identity access (least-privilege RBAC) to the Prod, NonProd/CVAL, and DR-hosting subscriptions (see T2). All mutations are attributed to the control identity and audited (GOV-002). [Decided]

5.2 CRG hierarchy and cross-subscription sharing (Section 25, FC-06)

- **Three-CRG model per region:** exactly one Prod, one Non-Prod/CVAL, and one DR-standby CRG target per region per scope (CAP-001). [Decided]
- **Zone alignment (FC-06):** cross-subscription CRG sharing requires a validated `ZoneMappingRecord` — logical zones differ across subscriptions,

so sharing is only valid when physical zones align (CAP-014..019).
[Decided]

- **Sharing limits:** provider→consumer sharing validated for SKU/region/zone/authorisation; a per-provider consumer ceiling applies. [Decided]

T3 — CRG hierarchy & cross-subscription sharing (zone alignment)

flowchart TB

```

    subgraph RegionR [Region R]
        ProdCRG[Prod CRG]
        NonProdCRG[NonProd/CVAL CRG]
        DRCRG[DR-standby CRG]
    end
    ProdCRG -->|reserves| ZmapP{{Zone map: logical->physical}}
    NonProdCRG --> ZmapN{{Zone map}}
    DRCRG --> ZmapD{{Zone map}}
    ZmapP --> ShareA[Consumer subscription A<br/>zone-aligned share]
    ZmapN --> ShareB[Consumer subscription B]
    ZmapD --> ShareC[DR consumer subscription]
    note1[FC-06: share valid only when physical zones align<br/>ZoneMappingRecord required]
    classDef n fill:#fff3cd,stroke:#e0a800;

```

5.3 Quota-group architecture — single-pool decision (Section 26, QUA-004)

Primary model: one governed quota group per region and quota family, covering Prod + NonProd/CVAL + DR together. All eligible default per-region quota is collected into the pool (QUA-003) and allocated on demand. Prod and DR are protected inside the shared pool by **logical earmarks** enforced by the Quota-Pool Manager at allocation/reclamation time — not by physical group separation. Because all three environments share one pool, an emergency DR draw needs **no cross-group transfer**. [Decided]

Exception topology: two (or more) groups per region, used **only** when Azure Quota Group limits or a mandatory Prod-isolation governance boundary make one pool impossible; each case is recorded in the Decision Log and reverts to one pool when the constraint lifts (DEP-001). [Decided]

The pool arithmetic (Pool_Limit, Prod_Reserved_Floor, DR_Earmark_vCPU, Allocatable_NonProd, Pool_Headroom) is specified in Section 8.3. See T4.

T4 — Quota-group / pool architecture (single-pool model)

flowchart TB

```

    subgraph Pool [Single governed quota pool per region/quota family]
        direction TB
        PF[Prod_Reserved_Floor<br/>= Prod_Used + Prod_Growth_Buffer<br/>never allocatable to]
        DE[DR_Earmark_vCPU<br/>= Destination_DR_Requirement x vCPU max-not-sum<br/>never all]
    end

```

```

        AN[Allocatable_NonProd<br/>= Pool_Limit - Prod_Reserved_Floor - DR_Earmark - NonProd
        PH[Pool_Headroom = Pool_Limit - Pool_Used]
    end
    Pool -->|emergency DR draw: pool headroom first,<br/>then reclaimable NonProd - no cross
    Exc[Exception topology only:<br/>Prod group + NonProd+DR group<br/>when Azure limit / Pr
    classDef e fill:#f8d7da,stroke:#dc3545;

```

6. Data Architecture (Section 34, DAT-001..006)

6.1 Store characteristics

A versioned document store with per-document optimistic concurrency (policy_version / _etag-style guards), freshness metadata on every snapshot, and append-only operation/audit records (DAT-001..006, NFR-007). Every decision-driving read carries a capacity_snapshot_ref for deterministic replay. [Decided]

6.2 Core entities

The three v2.2-critical entities are **CustomerSeedRecord**, **SourceDestinationDRIndex**, and **CVALEarmarkRecord**, alongside reservation/quota snapshots, PlacementPolicy, and operation/audit records.

```

CustomerSeedRecord {                                     # PLC-003/004/005 - first-placement authority
    customer_realm_id      : string (PK)
    geography              : string (PK)
    production_region      : string
    cval_region            : string
    dr_region              : string | "NOT_OFFERED"
    products_covered       : list<string>
    decision_timestamp      : datetime
    policy_version          : string
    capacity_snapshot_ref   : string
    exception_ref           : string | null
    approval_metadata       : ApprovalRecord
    input_mode              : "SPECIFIC_REGION" | "GEOGRAPHY_EXCEPTION"
}

SourceDestinationDRIndex {                               # DR-018 - reverse-of-seed, drives activation
    source_region           : string (PK)
    destination_region      : string (PK)
    customer_realm_id       : string (PK)
    standby_instance_set    : list<InstanceRef>
    sku_family              : string
    quantity_vcpu           : int

```

```

        activation_state      : STANDBY | ACTIVATING | ACTIVE | FAILBACK_PENDING
        last_updated          : datetime
        policy_version         : string
    }

    CVALEarmarkRecord {
        customer_realm_id      : string (PK)
        region                  : string (PK)
        cval_capacity_vcpu     : int
        dr_earmarked_vcpu      : int
        co_located              : bool
        last_updated           : datetime
    }
    # PLC-010 - prevents DR/CVAL double-count
    # portion counted toward DR, not live CVAL headro

```

Supporting entities: **RegionalSnapshot** (freshness-stamped capacity/quota/usage/zone data), **PlacementPolicy** (weights, buffers, thresholds, catalogue — versioned), **ReservationState/QuotaPoolState**, **OperationRecord** (saga + compensation), **AuditEvent** (append-only), **ActivationRecord** (DR-019, Section 9.3), **ZoneMappingRecord** (FC-06), **SharingRelationship**, **DRDistributionPlan**.

T5 — Data model / ERD (incl. seed record, SourceDestinationDRIndex, CVAL earmark)

```

erDiagram
    CustomerSeedRecord ||--o{ SourceDestinationDRIndex : "reverse view"
    CustomerSeedRecord ||--o{ CVALEarmarkRecord : "co-location"
    CustomerSeedRecord }o--|| PlacementPolicy : "policy_version"
    RegionalSnapshot ||--o{ QuotaPoolState : "captures"
    QuotaPoolState ||--o{ ReservationState : "governs"
    SourceDestinationDRIndex ||--o{ ActivationRecord : "activates"
    DRDistributionPlan ||--o{ SourceDestinationDRIndex : "aggregates"
    OperationRecord ||--o{ AuditEvent : "emits"
    CustomerSeedRecord {
        string customer_realm_id PK
        string geography PK
        string production_region
        string cval_region
        string dr_region
        string input_mode
        string policy_version
    }
    SourceDestinationDRIndex {
        string source_region PK
        string destination_region PK
        string customer_realm_id PK
        int quantity_vcpu
    }

```

```

        string activation_state
    }
    CVALEarmarkRecord {
        string customer_realm_id PK
        string region PK
        int cval_capacity_vcpu
        int dr_earmarked_vcpu
        bool co_located
    }
    QuotaPoolState {
        string region PK
        int pool_limit
        int prod_reserved_floor
        int dr_earmark_vcpu
        int allocatable_nonprod
    }
}

```

6.3 Freshness & versioning

Every snapshot carries `collected_at`; the readiness gate compares `now()` - `collected_at` against `max_snapshot_age` (RDY-004). Policy and scope files are versioned and decision-traceable; the version that drove any decision is persisted with it (DAT-005, GOV-003). [Decided]

7. State Model & Concurrency (Section 29, NFR-002/007)

7.1 Engine mode machine (five states)

STEADY_STATE → DR_DECLARATION_PENDING → DR_EVENT_ACTIVE → FAILBACK_PENDING → STEADY_STATE
INCIDENT_HOLD (safe degraded)

Steady-state reconciliation and auto-increase run **only** in STEADY_STATE; emergency transfer is callable **only** in DR_EVENT_ACTIVE. Entering DR_EVENT_ACTIVE does not auto-authorise service-impacting CVAL action — each action retains its own policy gate. INCIDENT_HOLD is the safe degraded state entered when inputs cannot be trusted. [Decided]

7.2 Concurrency & reservation-of-intent

- **Optimistic concurrency:** every mutating operation reads an expected version and fails the write on mismatch, forcing a re-read (NFR-007). [Decided]
- **Reservation-of-intent:** placement writes an atomic hold before the seed record so two concurrent onboardings cannot double-commit the same capacity (NFR-002). [Decided]

- **Saga + compensation:** multi-step mutations use `OperationRecord` with a compensation chain for rollback; `VM_ImpactRecord` is immutable for any VM state change. [Decided]

T6 — Five-state engine machine

stateDiagram-v2

```

[*] --> STEADY_STATE
STEADY_STATE --> DR_DECLARATION_PENDING: DR declared (authorised)
DR_DECLARATION_PENDING --> DR_EVENT_ACTIVE: activation approved
DR_DECLARATION_PENDING --> STEADY_STATE: declaration withdrawn
DR_EVENT_ACTIVE --> FAILBACK_PENDING: failback authorised
FAILBACK_PENDING --> STEADY_STATE: failback complete (reversible)
STEADY_STATE --> INCIDENT_HOLD: inputs untrusted / stale
DR_EVENT_ACTIVE --> INCIDENT_HOLD: safety trip
INCIDENT_HOLD --> STEADY_STATE: inputs restored

```

8. Placement Scoring & Forecasting (Section 28)

8.1 Input modes and pipeline

The default input is an **exact production region** which ACRME **validates** (it does not derive). A **geography** input is an **exception path** requiring approval + customer acknowledgement, after which the derived region becomes the fixed seed (ADR-001, PLC-001/002). Restricted regions never enter scoring — they route to the exception workflow.

T7 — Staged placement pipeline

flowchart LR

```

In[Placement request] --> HC[Hard constraints HC-1..HC-10]
HC -->|fail| Rej[Reason state / exception]
HC -->|pass| Sel{Input mode}
Sel -->|exact region default| Val[Validate & accept]
Sel -->|geography exception| Score[Score Standard regions PS_Prod]
Score --> Pick[argmax -> fixed seed]
Val --> Seq[Sequential CVAL then DR selection]
Pick --> Seq
Seq --> Seed[Write seed + DR index]

```

T8 — Prod region input modes (exact vs geography exception)

flowchart TD

```

R[Region input] --> Q{Exact region or geography?}
Q -->|Exact Standard| V[Validate HC-1..HC-10 -> accept, no argmax]
Q -->|Exact Restricted| X[Exception Deployment Workflow]

```

```

Q -->|Geography| E{Exception approved + acknowledged?}
E -->|No| B[POLICY_BLOCKED]
E -->|Yes| D[argmax PS_Prod over Standard regions -> fixed seed]

```

8.2 Scoring formulas (weights =0.30/ =0.20/ =0.25/ =0.15/ =0.10)

All five weights are tunable `PlacementPolicy` constants summing to 1.0; every component is clamped `Clamp(x)=max(0,min(1,x))` before weighting. [Assumed]

PlacementPolicy weight-sum validation gate [Decided] The weight-sum invariant (+ + + + = 1.0) is a **mathematical requirement**, not a convention. Without it, placement scores fall outside [0, 1], rendering threshold gates and cross-version comparisons invalid regardless of which region `argmax` selects.

Responsibility: Config / Scope-File Service. This service is the sole gatekeeper for `PlacementPolicy` versions. The validation gate runs at **policy version load time**, before the version is published to the Placement & Scoring Engine. The Placement Engine never receives a policy version that has not passed this gate.

Formal gate logic:

```

weight_sum = policy.alpha + policy.beta + policy.gamma + policy.delta + policy.epsilon
tolerance = 0.001 # IEEE 754 floating-point rounding only; not a margin for operator error

```

```

IF |weight_sum - 1.0| > tolerance:
    REJECT version - do NOT publish
    RAISE PolicyValidationError(code="WEIGHT_SUM_INVARIANT_VIOLATED", computed_sum=weight_sum)
    WRITE AuditEvent(event_type="POLICY_VALIDATION_FAILED", severity="ERROR", policy_version=version)
ELSE:
    PUBLISH version to Placement Engine
    WRITE AuditEvent(event_type="POLICY_VERSION_PUBLISHED", weight_sum=weight_sum)

```

Consequences of a non-1.0 sum:

Weights sum to	Scoring effect	Operational consequence
> 1.0 (e.g. 1.10)	Best possible score > 1.0 — scores inflated	Placement thresholds pass regions that should be blocked
< 1.0 (e.g. 0.90)	Best possible score < 1.0 — scores compressed	Placement thresholds block regions that should pass; cross-version comparisons invalid

Tuning constraint: when any weight is increased, the remaining weights must decrease by the same total. The tuning produces a valid new sum before the updated policy version is submitted:

Valid: =0.40, =0.15, =0.25, =0.10, =0.10 → sum = 1.00
Invalid: =0.40, =0.20, =0.25, =0.15, =0.10 → sum = 1.10 -
gate rejects

Test catalogue entries (T-WV-01..T-WV-05): see Calculation Logic Reference Section “Policy weight-sum validation rule” for the full test matrix covering the valid case, both invalid-sum directions, IEEE 754 rounding tolerance, and the tolerance boundary. [Decided]

```
PS_Prod(r)      = 0.30·Clamp(nonprod.effective_free / prod.quantity)
                + 0.20·Clamp(prod.quota_headroom / prod.quota_limit)
                + 0.25·Clamp(1 - prod_customer_count / total_customers)
                + 0.15·Clamp(dr.coverage_ratio)
                + 0.10·Clamp(az_count / 3)

PS_NonProd(r)   = 0.30·Clamp(nonprod.effective_free / nonprod.quantity)
                + 0.20·Clamp(nonprod.quota_headroom / nonprod.quota_limit)
                + 0.25·Clamp(1 - nonprod_customer_count / total_customers)
                + 0.15·Clamp(nonprod.effective_free / nonprod.quantity)
                + 0.10·Clamp(az_count / 3)

PS_DR(r)        = 0.30·Clamp(dr.free_slots / dr.quantity)
                + 0.20·Clamp(dr.quota_headroom / dr.quota_limit)
                + 0.25·Clamp(1 - dr_customer_count / total_customers)
                + 0.15·min(1.0, dr.coverage_ratio / dr_ratio_target)
                + 0.10·Clamp(az_count / 3)
```

PS_Prod is dual-purpose (derive on exception; validate/audit otherwise). Every candidate score + policy version + snapshot ref is written to the `OperationRecord` for deterministic replay. A reviewer-recommended pilot PS_NonProd variant (removing the / duplication) is recorded in the Calculation Logic Reference; the formula above is the approved design-of-record. Scoring runs in shadow/recommendation mode until empirically validated. [Decided]

8.3 Quota-pool arithmetic (single-pool, Section 26/QUA-004)

```
Pool_Limit(region) = Prod_CRG_qty·vCPU·(1+prod_growth_buffer)
                  + NonProd_CRG_qty·vCPU·(1+nonprod_growth_buffer)
                  + DR_Earmark_vCPU(region)
                  + emergency_transfer_headroom_vcpu
```

```
Prod_Reserved_Floor = Prod_Used_vCPU + Prod_Growth_Buffer_vCPU      # never allocatable t
DR_Earmark_vCPU      = Destination_DR_Requirement(region)·vCPU      # max-not-sum, never
```

```

Allocatable_NonProd = Pool_Limit - Prod_Reserved_Floor - DR_Earmark_vCPU - NonProd_Used_vCPU
Pool_Headroom       = Pool_Limit - Pool_Used
Emergency_DR_Available = Pool_Headroom + reclaimable_NonProd_above_committed

```

NonProd allocation fail-safes when `Allocatable_NonProd` reaches zero. A dual-validation detector recomputes the DR earmark from authoritative assignment data; disagreement disables automatic NonProd expansion until reconciled (Scenario 9). [Decided]

8.4 Forecasting vs reconciliation floor

```

Target_Reserved_Capacity = Allocated_VM_Count + Configured_Buffer # CAP-003 continued
Forecast_Quantity        = ceil(Forecast_Peak * (1 + Growth_Buffer) + DR_Buffer) # ADR-004 proposed

```

The reconciliation floor is enforced every cycle; the forecast is advisory (horizon {30,60,90} days) and raises `ForecastApproachingQuotaLimit` with 14-day lead time at 80% of quota limit. Auto-increase target = max(floor, forecast). [Decided]

9. DR Sizing & Activation Algorithms (Section 31, App. A.6/A.7/A.8, App. D)

9.1 Max-not-sum sizing (DR-017, ADR-005)

```

Destination_DR_Requirement(d) = MAX over each non-concurrent source s protected by d
                               ( Workload_Portion(s -> d) )
DR_Capacity_Gap(d)            = max(0, Destination_DR_Requirement(d) - Usable_Destination_Capacity(d))
Overcommit_Ratio(d)           = SUM(Workload_Portion(s->d)) / MAX(Workload_Portion(s->d))

```

Single-failure assumption (DR-001): one production region fails at a time, so destination `d` needs standby for its **largest** protected source only; standby is shared/overcommitted across mutually exclusive events. `Usable_Destination_Capacity(d)` may include approved bootstrap headroom, available CRG reservations, releasable CVAL (after earmark check), and approved sharing/expansion. A per-scope **SUM override (C-11)** covers contractual simultaneous-failure protection with no code change. [Decided]

DR_NOT_OFFERED geographies are excluded from this sizing (DR-014, DEC-001). Sources in a geography flagged `DR_NOT_OFFERED` — the Middle East today — contribute **no** `Workload_Portion(s → d)` term, produce **no** `SourceDestinationDRIndex` entry, and reserve **no** `DR_Earmark_vCPU` at any destination. Their production is still sized and governed, but they carry no DR requirement, so they never enter `Destination_DR_Requirement`, `DR_Capacity_Gap`, or `Overcommit_Ratio`. Should Legal clear DEC-001 (config flip `dr_not_offered["Middle East"] = false`, no code change), the

affected sources begin contributing their max-not-sum term at that point. [Documented]

T12 — Max-not-sum sizing (overcommit ratio)

flowchart TB

S1[Source R1 -> R2: 120 cores]

S3[Source R3 -> R2: 80 cores]

S1 --> R2

S3 --> R2

R2[Destination R2 standby
SUM = 200 cores over-provisions
MAX = 120 cores sized]

classDef d fill:#d4edda,stroke:#28a745;

9.2 Standby activation waves & staged acquisition (DR-006, DR-019)

On authorised declaration the DR Orchestrator reads the `SourceDestinationDRIndex` for the failed source and activates customers in business-priority waves (P0 platform → P-1 highest → P-N), transitioning each `associated` → `allocated`. Per wave, staged acquisition runs in order:

Stage 1: approved bootstrap / pre-staged headroom (zero Azure wait)

Stage 2: available reservation + quota already in destination

Stage 3: shut down / disassociate eligible CVAL (after DR-010 authorised trigger)

Stage 4: share/reassign reservations within supported region/zone boundaries

Stage 5: allocate pooled quota to DR subscriptions

Stage 6: request additional Azure quota/capacity where required

Stage 7: report unrecoverable capacity gaps

Stages 1–2 start immediately from the single pool’s headroom, so P0/P-1 recovery does not wait on throttled Azure deallocation APIs. [Decided]

9.3 Activation state tracking & fallback (DR-013)

```
ActivationRecord {
    customer_realm_id, source_region, destination_region,
    priority_wave : int,
    activation_state : PENDING | ACTIVATING | ACTIVE | FAILED | FAILBACK_PENDING,
    stage_reached : int (1..7), capacity_gap_vcpu : int,
    activated_at, approved_by, incident_id
}
```

Fallback reverses `ACTIVE` → `FAILBACK_PENDING` → `STANDBY` in reverse priority order; seed record and DR index are preserved (no re-derivation). [Decided]

T11 — DR standby activation sequence (index lookup → waves → staged acquisition)

sequenceDiagram

```
participant DRC as DR coordinator
participant DRO as DR Orchestrator
participant IDX as SourceDestinationDRIndex
participant QPM as Quota-Pool Manager
participant AZ as Azure
DRC->>DRO: Declare source-region failure (authorised)
DRO->>DRO: engine_mode = DR_EVENT_ACTIVE
DRO->>IDX: Query standby sets for failed source
IDX-->>DRO: Customers + priority waves
loop each priority wave (P0 -> P-1 -> P-N)
    DRO->>QPM: Stage1-2 draw pool headroom + reservation
    QPM-->>DRO: Earmark released (no cross-group transfer)
    alt deficit remains
        DRO->>QPM: Stage3 CVAL release / Stage4 share / Stage5 pooled quota
        DRO->>AZ: Stage6 request additional capacity/quota
        AZ-->>DRO: Granted or deficit
    end
    end
    DRO->>DRO: associated -> allocated (ActivationRecord)
    DRO->>DRC: Wave complete + audit checkpoint
end
DRC->>DRO: Failback authorised
DRO->>DRO: ACTIVE -> FAILBACK_PENDING -> STANDBY (reverse order)
```

9.4 Tier escalation (Section 32) — see Section 13 baseline

Emergency capacity during DR_EVENT_ACTIVE escalates Tier 1 (direct expansion) → Tier 2 (quota-neutral transfer within the shared pool) → Tier 3 (destructive, dual-approval, **blocked in Phase 1**). In the single-pool model Tier 2 is intrinsically quota-neutral (draw and release hit the same pool). [Decided]

T10 — Three-tier emergency transfer escalation

flowchart TD

```
Ev[DR event active] --> T1{Tier 1 headroom available?}
T1 -->|Yes| A1[DirectExpansion - automated additive]
T1 -->|No| T2{Tier 2 approved?}
T2 -->|Yes| A2[QuotaNeutralTransfer - reduce NonProd, expand DR in same pool]
T2 -->|No| T3{Tier 3 allowed?}
T3 -->|Phase 1: BLOCKED| M[Manual escalation]
T3 -->|Later phase: dual-approval| A3[DestructiveTransfer - changes VM associations]
```

T13 — Distributed DR reference model (Section 12A)

Reused reference diagram (ADR-005 Section 12A): `adr/diagrams/acrme_three_region_capacity_model.png`
— each region concurrently hosts Prod + CVAL + DR standby tagged with its `src Rn` source; each destination is sized by max-not-sum against its largest single protected source.

10. Steady-State Capacity Lifecycle (Section 30)

The normative 10-step sequence runs only in `STEADY_STATE`:

1. Detect threshold crossing (utilisation auto-increase threshold).
2. Re-read current CR, quota, sharing, assignment state.
3. Create `CapacityIncreaseRequest`.
4. Calculate `target = max(Allocated+Buffer floor, Forecast_Quantity)`.
5. Require operator approval (Phase 1).
6. Submit quota action only if validated as required.
7. Wait for confirmed quota state - no assumed propagation SLA.
8. Update CR quantity.
9. Confirm actual quantity.
10. Refresh snapshot and close the request.

Triggers: `dr_autoincrease_threshold=0.35`, `prod/nonprod_autoincrease_threshold=0.20`,
debounce cooldown 30 min per (region+CRG type). Guarded scale-down never
drops below `Allocated`, within a DR window, during maintenance exclusion,
or against cost policy; reductions go to zero, never deletion (`CAP-009/010`).
[Decided]

T9 — Steady-state 10-step capacity lifecycle

sequenceDiagram

```
participant REC as State Reconciler
participant OP as Operator
participant QPM as Quota-Pool Manager
participant AZ as Azure
REC->>REC: 1. Detect threshold crossing
REC->>AZ: 2. Re-read CR/quota/sharing/assignment
REC->>REC: 3. Create CapacityIncreaseRequest
REC->>REC: 4. Target = max(Allocated+Buffer, Forecast)
REC->>OP: 5. Require approval (Phase 1)
OP-->>REC: Approved
REC->>QPM: 6. Submit quota action if required
QPM->>AZ: 7. Await confirmed quota state
AZ-->>QPM: Confirmed
QPM->>AZ: 8. Update CR quantity
```

AZ-->>QPM: 9. Confirm actual quantity
 REC->>REC: 10. Refresh snapshot + close request

11. API Architecture (Section 35, INT-001..007)

11.1 Contracts

POST /readiness → DeploymentReadinessResult (query; idempotent)
 in : { customer_realm_id, region (exact), sku_family, environment, product, intended_demand
 out: { readiness_state, reasons[], operation_handle, snapshot_ref, policy_version }

POST /placement → PlacementResult (idempotent; writes seed + DR index)
 POST /capacity/increase → OperationRecord (operator-gated; idempotent)
 POST /dr/declare → OperationRecord (authorised; flips engine_mode)
 POST /dr/activate-wave → ActivationRecord[] (per priority wave)
 POST /quota/allocate → OperationRecord (single-pool allocation)
 POST /sharing/activate → OperationRecord (zone-aligned share)
 GET /audit, GET /state → read-only

11.2 Idempotency, versioning, error/readiness codes

- Every mutating call carries an **idempotency key** and an **expected-state version**; retried calls return the original result, mismatched versions return a concurrency error (INT-004..006, NFR-007). [Decided]
- Readiness states are the machine-readable contract to AEP:

READY | READY_WITH_RISK | QUOTA_DEFICIT | RESERVATION_DEFICIT |
 CAPACITY_UNAVAILABLE | STALE_STATE | POLICY_BLOCKED | VALIDATION_REQUIRED

- AEP **polls** operation handles rather than assuming synchronous completion, and never treats provider quota as proof of consumer deployability (INT-007, POC-001). [Assumed]

11.3 Readiness gate logic (RDY-001)

READY requires ALL of:

1. target_region approved_region_catalogue (REG-001)
2. az_count 1 in target_region (CAP-011)
3. sku supported_skus(target_region)
4. reservation_policy known for (subscription, region, zone, sku)
5. reservation exists (if required by policy)
6. reserved_quantity requested_count OR approved_over_allocation active
7. consumer_subscription_quota requested_units (QUA-013 - POC-gated)
8. snapshot_age max_snapshot_age_seconds (RDY-004)
9. no blocking policy exception or hard-constraint violation

Quota_Deficit = max(0, Requested_Deployment_Units - Available_Quota)
 Available_Quota = Assigned_Regional_VM_Family_Quota - Current_Regional_VM_Family_Usage

12. Security (Section 36, GOV-001..009)

- **RBAC & managed identity:** the control identity holds least-privilege, scoped roles per target subscription; no standing user credentials for mutations (GOV-001/002). [Decided]
 - **Approval gates:** capacity/quota increase, DR declaration, CVAL release, Tier 2 transfer, seed change, and overrides are policy-gated with recorded approver identity (GOV-005..007). [Decided]
 - **Break-glass:** an audited emergency-elevation path with dual control and time-boxed elevated RBAC; every break-glass action is immutably logged (GOV-008). [Decided]
 - **Secrets:** managed identity + platform secret store; no secrets in config or code. [Decided]
 - **Audit:** append-only `AuditEvent` for every decision and mutation with before/after values and correlation IDs (GOV-004/009, DAT-005). [Decided]
-

13. Observability Implementation (Section 37, OBS-001..005)

Signal	Metric / alert	Requirement
Buffer deficit	<code>reservation_buffer_deficit</code> gauge; alert when reserved < Allocated+Buffer	OBS-001, CAP-003
Per-destination DR coverage	<code>dr_max_source_coverage</code> alert when usable < Destination_DR_Requirement(d) (the MAX, not the SUM)	OBS-002, DR-017
DR capacity gap	<code>dr_capacity_gap(d)</code> gauge	OBS-002
Source destination mapping	dashboard of the <code>SourceDestinationDRIndex</code>	OBS-004
Overcommit ratio	<code>overcommit_ratio(d)</code> ; alert above safety ceiling	FIN-006..008, POC-011
Activation state	per-wave <code>ActivationRecord</code> progress	OBS-003, DR-019

Signal	Metric / alert	Requirement
Idle cost	idle reservation cost surfaced	OBS-005, FIN-001/002
Staleness	snapshot_age ; alert/refresh on breach	RDY-004, NFR-002
Forecast lead-time	ForecastApproachingQuotaLimit (14-day lead at 80%)	OBS-001, ADR-004

All metrics carry region/environment/scope dimensions and policy version for replay. [Decided]

14. NFRs & Resilience (Section 38, NFR-001..010)

- **Throttling resilience:** adaptive throttle manager initialises with ARM baselines (250 reads/5 min, 1,200 writes/hour per subscription) and adapts to observed **429/Retry-After**; reconciliation cadence and estate size are bounded accordingly (Section 14 calc). [Documented]
- **Estate sizing:** $CRG_total = R \cdot Eclass \cdot Z \cdot SKUset \cdot IsolationFactor$; **Requests_per_cycle** summed across read classes / 300 s cycle to stay within budget. [Derived]
- **Degraded mode:** on untrusted/stale inputs the engine enters **INCIDENT_HOLD** / returns **STALE_STATE** and fails safe rather than driving placement from stale data (NFR-002). [Decided]
- **Simulation / dry-run:** placement, reconciliation, and **DR failover** can run in dry-run/simulation, producing the would-be decisions and readiness states without mutating Azure (NFR-006). [Decided]
- **Idempotency & concurrency:** all mutations idempotent + optimistic concurrency (NFR-007). Determinism/replay from **OperationRecord**. [Decided]

15. Integration (Section 33, Section 6 review)

- **AKS / VMSS:** reservations back node pools / scale sets; ACRME readiness gates their capacity before scale-out; associations are tracked for reconciliation. [Decided]
- **AEP:** the sole provisioning consumer of readiness; interacts only through the idempotent readiness/placement contract (Section 11) and polls operation handles. [Decided]
- **Azure control planes:** Compute (CRGs/reservations), Reservations, and **groupQuotas** (single-pool). Single-pool enforcement depends on **groupQuotas / groupType** maturity (DEP-001). [Documented]

16. POC & Validation Dependencies (Section 42)

POC	Question	Gates	Status
POC-001	Does provider quota under a shared reservation guarantee consumer deployability?	Readiness gate step 7; INT-007	Top unknown
POC-006	Distributed DR topology behaves as modelled	ADR-005 Section 12A; DR-016	Required
POC-007	Lean bootstrap sizing is sufficient to initiate recovery	DR-007 bootstrap targets	Required
POC-011	max-not-sum overcommit is safe at platform scale	DR-017 sizing; overcommit safety ceiling; FIN-006..008	Required before production dependency
POC-005	VM state semantics (associated/allocated/deallocated)	DR-019 activation staging	Dependency
POC-031/032	Tier 2 quota-neutral transfer within a group/pool	Section 13 baseline Tier 2	Required

Production reliance on max-not-sum DR sizing and single-pool consumer-quota behaviour is **gated** on POC-011 and POC-001 respectively; until then, readiness validates consumer quota explicitly and DR sizing carries the SUM-override safety option. [Decided]

17. Traceability — Requirement/Deviation → Component/Algorithm/Entity

Requirement group (v2.2 IDs)	Component(s)	Algorithm / formula	Entity / diagram
REG-001..003	Config/Scope-File, Placement Engine	catalogue validation; input modes	PlacementPolicy ; T8
ENV-001..007	Placement Engine, DR Orchestrator	env separation; role flip	CustomerSeedRecord ; T7
CAP-001..019	Inventory Collector, State Reconciler	Allocated+Buffer floor; zero-not-delete; over-alloc	ReservationState ; T9, F5
QUA-001..014	Quota-Pool Manager	Section 8.3 pool arithmetic; earmarks; quota-as-governor	QuotaPoolState ; T4
RDY-001..004	Readiness API	Section 11.3 gate logic; staleness; quota deficit	readiness states; T6
PLC-001..010	Placement & Scoring Engine	PS_Prod/NonProd/DR seed reuse; co-location guard	CustomerSeedRecord , CVALEarmarkRecord ; T7, T8
DR-001..019	DR Orchestrator, Emergency Transfer	max-not-sum; gap; overcommit; activation waves; staged acquisition; tiers	SourceDestinationDRIndex , ActivationRecord ; T11, T12, T13, T10
FIN-001..008	Observability Emitter, Forecast Service	idle cost; overcommit ratio; cost-before-expansion	metrics; T12
INT-001..007	Readiness API	idempotency keys; expected-version; polling	OperationRecord ; T1
DAT-001..006	State Store, Config Service	freshness metadata; versioning; audit	all entities; T5
OBS-001..005	Observability Emitter	max-source coverage; source dest map; activation progress	dashboards; Section 13

Requirement group (v2.2 IDs)	Component(s)	Algorithm / formula	Entity / diagram
GOV-001..009	Governance & Audit	RBAC; approval gates; break-glass; immutable audit	AuditEvent ; Section 12
NFR-001..010	(cross-cutting)	throttle resilience; degraded mode; idempotency; simulation	OperationRecord ; Section 14
OPS-001..005	(runbooks)	DR declaration/activation/failback; quota allocation; sharing; override	ActivationRecord ; Section 10, Section 16
POC-001..011	(gating)	consumer quota; DR topology; bootstrap; overcommit safety	Section 16
DEC-001..003 / DEP-001	Config/Scope-File, Quota-Pool Manager	Middle East DR policy — current position DR_NOT_OFFERED, pending legal review (production allowed, no DR, Switzerland North path inactive until config flip); failback duration; geo-exception approver; groupQuotas maturity	PlacementPolicy ; Section 2.2, Section 5.3

Every Baseline v2.2 requirement group resolves to at least one component, algorithm, and entity above. Functional-level traceability is in the FDD Section 9.

18. Class Diagram (updated with new entities)

T14 — Domain class model (updated: +CustomerSeedRecord, +SourceDestinationDRIndex, +CVALEarmarkRecord; single-pool QuotaPoolState; readiness enum)

```
classDiagram
    class CustomerSeedRecord {
        +string customer_realm_id
        +string geography
        +string production_region
        +string cval_region
        +string dr_region
        +string input_mode
        +string policy_version
    }
    class SourceDestinationDRIndex {
        +string source_region
        +string destination_region
        +string customer_realm_id
        +int quantity_vcpu
        +ActivationState activation_state
    }
    class CVALEarmarkRecord {
        +int cval_capacity_vcpu
        +int dr_earmarked_vcpu
        +bool co_located
    }
    class QuotaPoolState {
        +int pool_limit
        +int prod_reserved_floor
        +int dr_earmark_vcpu
        +int allocatable_nonprod
        +int pool_headroom
    }
    class PlacementPolicy {
        +float alpha_beta_gamma_delta_epsilon
        +float prod_growth_buffer
        +float nonprod_growth_buffer
        +string version
    }
    class ActivationRecord {
        +int priority_wave
        +int stage_reached
        +int capacity_gap_vcpu
        +ActivationState state
    }
```

```

}
class ReservationState {
    +int reserved_quantity
    +int allocated_count
}
class OperationRecord {
    +string operation_handle
    +string expected_version
}
CustomerSeedRecord "1" --> "0..*" SourceDestinationDRIndex : reverse view
CustomerSeedRecord "1" --> "0..*" CVALEarmarkRecord
CustomerSeedRecord --> PlacementPolicy : policy_version
SourceDestinationDRIndex "1" --> "0..*" ActivationRecord
QuotaPoolState "1" --> "0..*" ReservationState
OperationRecord --> ReservationState : mutates

```

Note: the earlier UML summary's "two quota groups per region" is superseded by the **single governed QuotaPoolState** with logical earmarks (exception two-group topology only). [Decided]

Document Status: Draft for review · **Next:** ratify alongside the FDD; supersede the Production Readiness Review (archived). Production reliance on max-not-sum DR and single-pool consumer quota is gated on POC-011 / POC-001.